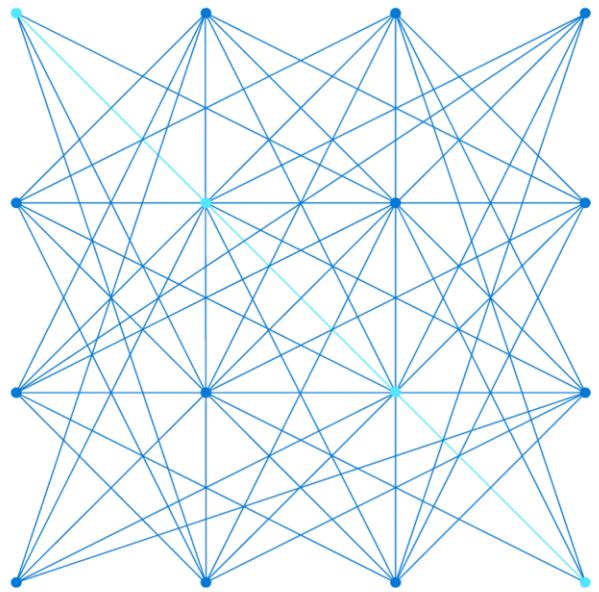




Implementing an Azure Data Solution [DP-200]



1

Agenda



About this course



Audience



Course agenda



Prerequisites

2

About this course

In this course, the students will implement various data platform technologies into solutions that are in line with business and technical requirements including on-premises, cloud, and hybrid data scenarios incorporating both relational and No-SQL data. They will also learn how to process data using a range of technologies and languages for both streaming and batch data

The students will also explore how to implement data security including authentication, authorization, data policies and standards. They will also define and implement data solution monitoring for both the data storage and data processing activities. Finally, they will manage and troubleshoot Azure data solutions which includes the optimization and disaster recovery of big data, batch processing and streaming data solutions

3

Course agenda

Module 1

Azure for the Data Engineer

Lesson 01 – Explain the evolving world of data

Lesson 02 – Survey the services in the Azure Data Platform

Lesson 03 – Identify the tasks that are performed by a Data Engineer

Lesson 04 – Describe the use cases for the cloud in a case study

Module 2

Working with Data Storage

Lesson 01 – Choose a data storage approach in Azure

Lesson 02 – Create an Azure Storage Account

Lesson 03 – Explain Azure Data Lake Storage

Lesson 04 – Upload data into Azure Data Lake

4

Course agenda (*continued #1*)

Module 3

Enabling team based Data Science with Azure Databricks

Lesson 01 – Explain Azure Databricks

Lesson 02 – Work with Azure Databricks

Lesson 03 – Read data with Azure Databricks

Lesson 04 – Perform transformations with Azure Databricks

Module 4

Building globally distributed databases with Cosmos DB

Lesson 01 – Create an Azure Cosmos DB database built to scale

Lesson 02 – Insert and query data in your Azure Cosmos DB database

Lesson 03 – Build a .NET Core app for Azure Cosmos DB in Visual Studio Code

Lesson 04 – Distribute your data globally with Azure Cosmos DB

5

Course agenda (*continued #2*)

Module 5

Working with relational data stores in the cloud

Lesson 01 – Explain SQL Database

Lesson 02 – Explain SQL Data Warehouse

Lesson 03 – Provision and load data in Azure SQL Data Warehouse

Lesson 04 – Import data into Azure SQL Data Warehouse using PolyBase

Module 6

Performing real-time analytics with Stream Analytics

Lesson 01 – Explain data streams and event processing

Lesson 02 – Data ingestion with Event Hubs

Lesson 03 – Processing data with Stream Analytics jobs

6

Course agenda (*continued #3*)

Module 7

Orchestrating data movement with Azure Data Factory

Lesson 01 – Explain how Azure Data Factory works

Lesson 02 – Create linked services and datasets

Lesson 03 – Create pipelines and activities

Lesson 04 – Azure Data Factory pipeline execution and triggers

Module 8

Securing Azure Data Platforms

Lesson 01 – Introduction to security

Lesson 02 – Key security components

Lesson 03 – Securing storage accounts and Data Lake Storage

Lesson 04 – Security data stores

Lesson 05 – Securing streaming data

7

Course agenda (*continued #4*)

Module 9

Monitoring and troubleshooting Data Storage and processing

Lesson 01 – Explain the monitoring capabilities that are available

Lesson 02 – Troubleshoot common data storage issues

Lesson 03 – Troubleshoot common data processing issues

Lesson 04 – Manage disaster recovery

8

Audience

Primary audience:

The audience for this course are data professionals, data architects, and business intelligence professionals who want to learn about the data platform technologies that exist on Microsoft Azure

Secondary audience:

The secondary audience for this course are individuals who develop applications that deliver content from the data platform technologies that exist on Microsoft Azure

9



10

Prerequisites

In addition to their professional experience, students who take this training should have technical knowledge equivalent to the following courses:

[Azure fundamentals](#)

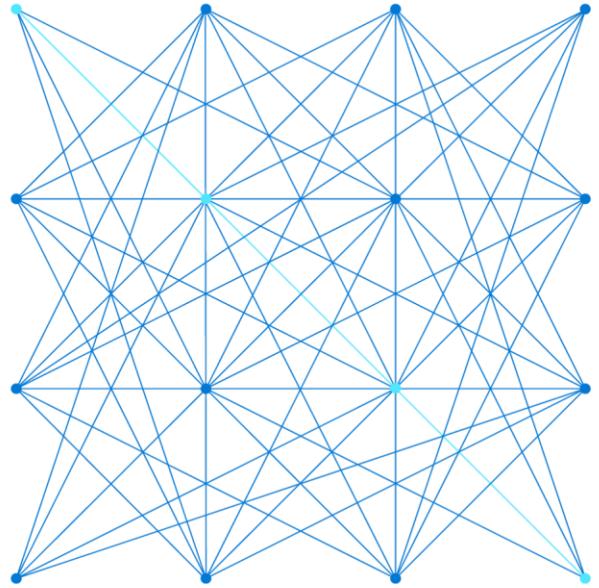


© Copyright Microsoft Corporation. All rights reserved.

11



Module 01: Azure for the Data Engineer



12

Agenda



Lesson 01 – Explain the evolving world of data



Lesson 02 – Survey the services in the Azure Data Platform



Lesson 03 – Identify the tasks that are performed by a Data Engineer



Lesson 04 – Describe the use cases for the cloud in a case study

13

Lesson 01: The evolving world of data



14

Lesson objectives



Data abundance



Differences between on-premises and cloud data technologies



How the role of the data professional is changing in organizations



Identify use cases impacted by these changes

15

Data abundance

Processes	Businesses are tasked to store, interpret, manage, transform, process, aggregate and report on data
Consumers	There are a wider range of consumers using different types of devices to consume or generate data
Variety	There's a wider variety of data types that need to be processed and stored
Responsibilities	A data engineer's role is responsible for more data types and technologies
Technologies	Microsoft Azure provides a wide set of tools and technologies

16

On-premises versus cloud technologies



Computing Environment



Licensing Model



Maintainability



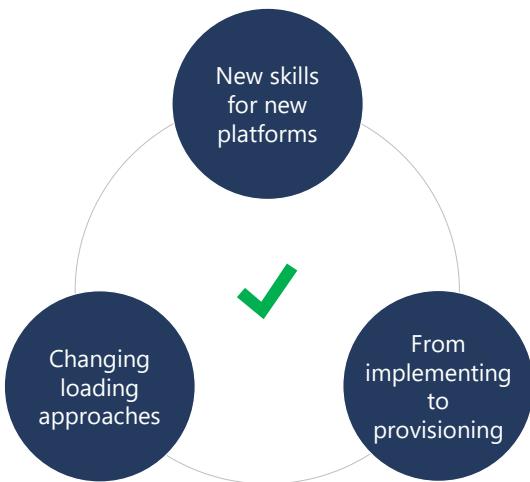
Scalability



Availability

17

Data engineering job responsibilities



18

Use cases for the cloud

Here are some examples of industries making use of the cloud

Web retail	Using Azure Cosmos DB's multi-master replication model along with Microsoft's performance commitments, Data Engineers can implement a data architecture to support web and mobile applications that achieve less than a 10-ms response time anywhere in the world
Healthcare	Azure Databricks can be used to accelerate big data analytics and artificial intelligence (AI) solutions. Within the healthcare industry, it can be used to perform genome studies or pharmacy sales forecasting at petabyte scale
IoT scenarios	Hundreds of thousands of devices have been designed and sold to generate sensor data known as Internet of Things (IoT) devices. Using technologies like Azure IoT Hub, Data Engineers can easily design a data solution architecture that captures real-time data

19

Lesson 02: Survey the services in the Azure Data Platform



20

Lesson objectives



The differences between structured and unstructured data



Azure Storage



Azure Data Lake Storage



Azure Databricks



Azure Cosmos DB



Azure SQL Database



Azure SQL Data Warehouse



Azure Stream Analytics



Additional Azure Data Platform Services

21

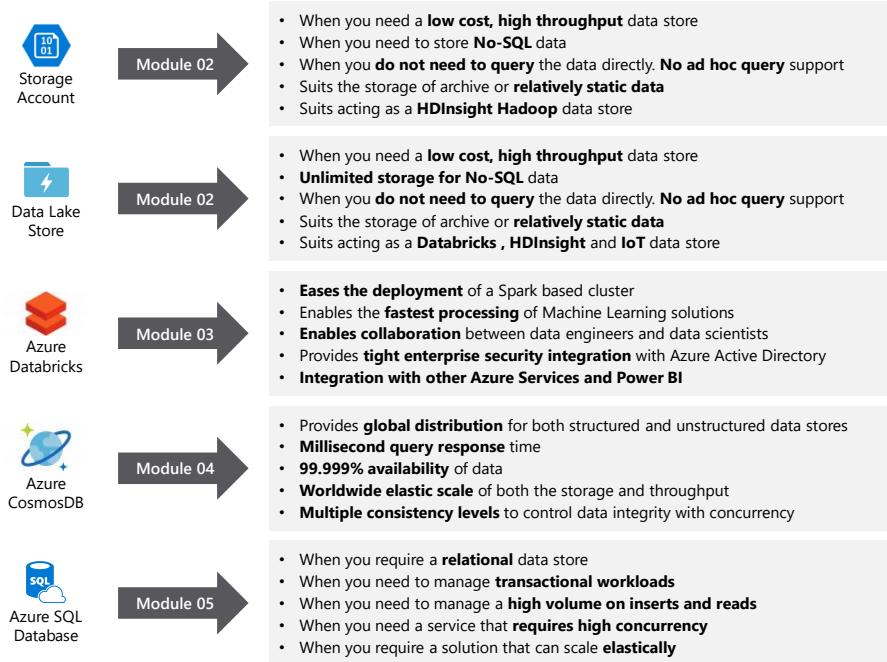
Structured versus Unstructured data

There are three broad types of data and Microsoft Azure provides many data platform technologies to meet the needs of the wide varieties of data

Structured	Semi-Structured	Unstructured
Structured data is data that adheres to a schema, so all of the data has the same fields or properties. Structured data can be stored in a database table with rows and columns	Semi-structured data doesn't fit neatly into tables, rows, and columns. Instead, semi-structured data uses _tags_ or _keys_ that organize and provide a hierarchy for the data	Unstructured data encompasses data that has no designated structure to it. Known as No-SQL, there are four types of No-SQL databases: <ul style="list-style-type: none"> • Key Value Store • Document Database • Graph Databases • Column Base

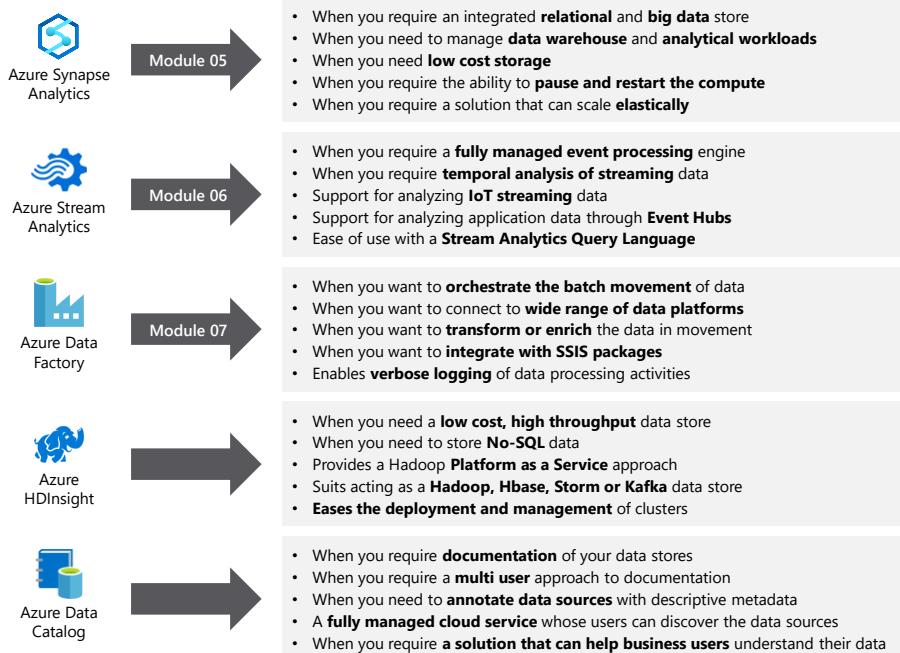
22

What to use for Data



23

What to use for Data



24

Lesson 03: Identify the tasks performed by a Data Engineer



25

Lesson objectives



List the new roles of modern data projects



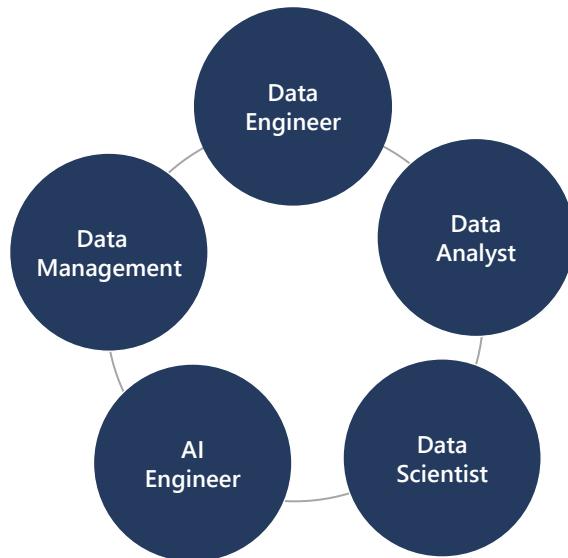
Outline data engineering practices



Explore the high-level process for architecting a data engineering project

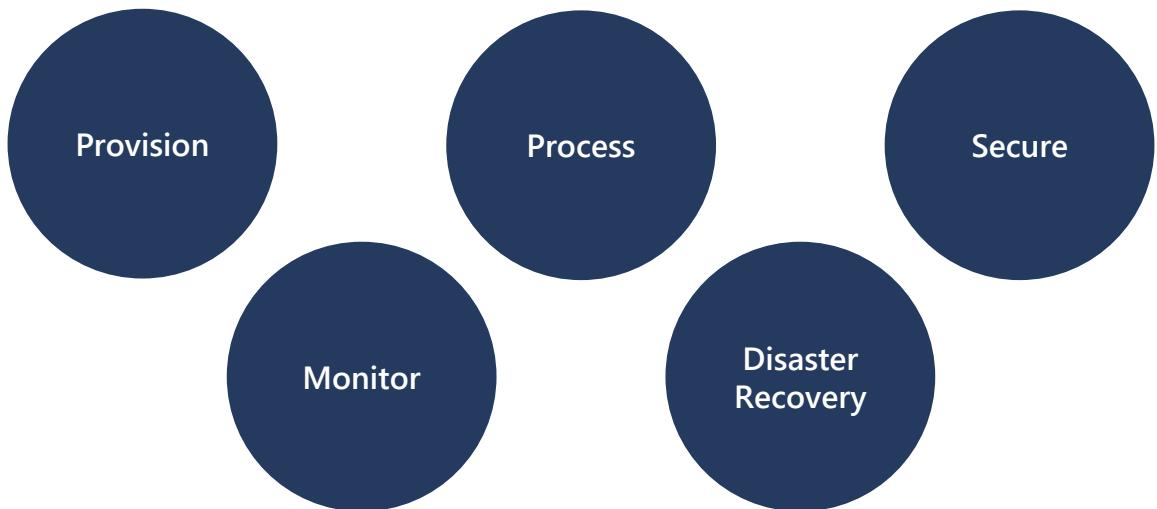
26

Roles in data projects



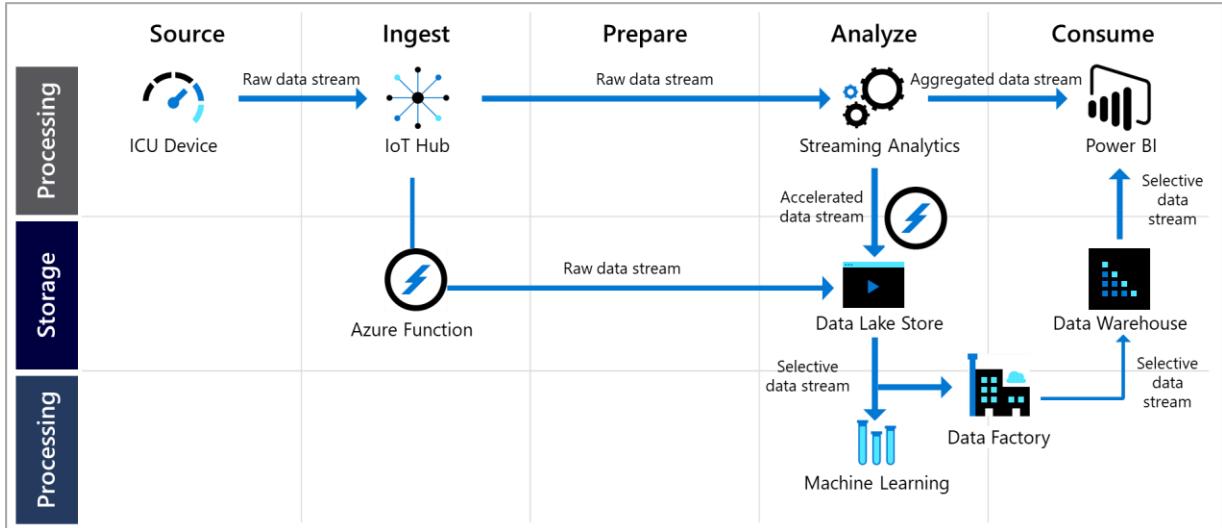
27

Data engineering practices



28

Architecting projects – An example

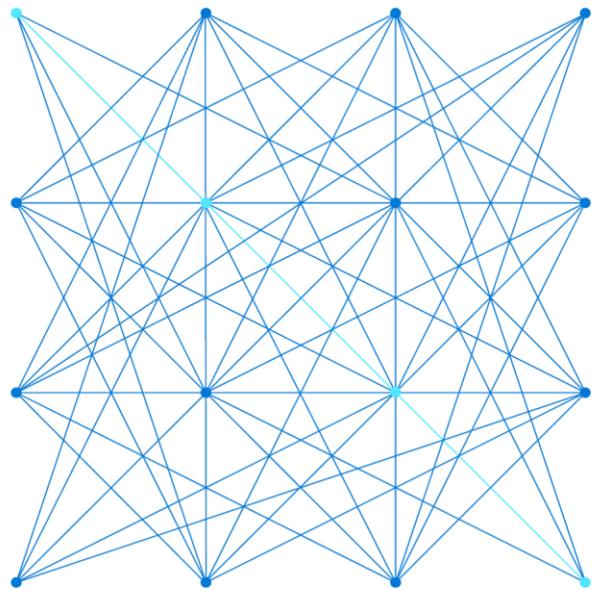


29



30

Module 02: Working with data storage



31

Agenda



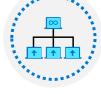
Lesson 01: Choose a data storage approach in Azure



Lesson 02: Create an Azure Storage Account



Lesson 03: Explain Azure Data Lake Storage



Lesson 04: Upload data into Azure Data Lake Store

32

Lesson 01: Choose a data storage approach in Azure



33

Lesson objectives



The Benefits of using Azure to store data



Compare Azure data storage with on-premises storage

34

Benefits of using Azure to store data

	Automated backup		Support for data analytics
	Global replication		Storage tiers
	Encryption capabilities		Virtual disks
	Multiple data types		

35

Comparing Azure to on-premises storage

The term “on-premises” refers to the storage and maintenance of data on local hardware and servers

Cost effectiveness	Reliability	Storage types	Agility
On-premises storage requires up-front expenses. Azure data storage provides a pay-as-you-go pricing model	Azure data storage provides backup, load balancing, disaster recovery, and data replication to ensure safety and high availability. This capability requires significant investment with on-premises solutions	Azure data storage provides a variety of different storage options including distributed access and tiered storage	Azure data storage gives you the flexibility to create new services in minutes and allows you to change storage back-ends quickly

36

Lesson 02: Create Azure storage account



37

Lesson objectives



Describe storage accounts



Determine the appropriate settings for each storage account



Choose an account creation tool



Create a storage account using the Azure portal

38

Storage accounts

What is a storage account?

It is a container that groups a set of Azure Storage services. Only data services can be included in a storage account such as *Azure Blobs*, *Azure Files*, *Azure Queues*, and *Azure Tables*

How many do you need?

The number of storage accounts you need is typically determined by your data diversity, cost sensitivity, and tolerance for management overhead

The number of storage accounts you need is based on:

Data diversity:

Organizations often generate data that differs in where it is consumed and how sensitive it is

Cost sensitivity:

The settings you choose for the account do influence the cost of services, and the number of accounts you create

Management overhead:

Each storage account requires some time and attention from an administrator to create and maintain

39

Storage account settings

[Home](#) > New > Storage account > Create storage account

Create storage account

Basics Networking Advanced Tags Review + create

Azure Storage is a Microsoft-managed service providing cloud storage that is highly available, secure, durable, scalable, and redundant. Azure Storage includes Azure Blobs (objects), Azure Data Lake Storage Gen2, Azure Files, Azure Queues, and Azure Tables. The cost of your storage account depends on the usage and the options you choose below.

[Learn more about Azure storage accounts](#)

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription * dtestao

Resource group * Select existing... Create new

Instance details

The default deployment model is Resource Manager, which supports the latest Azure features. You may choose to deploy using the classic deployment model instead. [Choose classic deployment model](#)

Storage account name * dtestao

Location * (US) South Central US

Performance Standard Premium

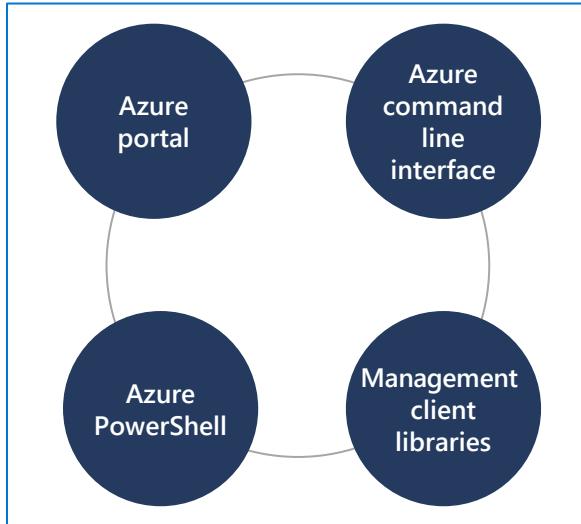
Account kind StorageV2 (general purpose v2)

Replication Read-access geo-redundant storage (RA-GRS)

Access tier (default) cool Hot

40

Storage account creation tool



41

Create a storage account



42

Lesson 03: Azure Data Lake Storage

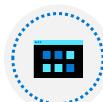


43

Lesson objectives



Explain Azure Data Lake Storage



Create an Azure Data Lake Store Gen 2 using the portal



Compare Azure Blob Storage and Data Lake Store Gen 2



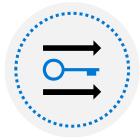
Explore the stages for processing Big Data using Azure Data Lake Store



Describe the use cases for Data lake Storage

44

Azure Data Lake Storage – Generation II



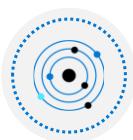
Hadoop access



Security



Performance



Redundancy

45

Create a Azure Data Lake Store (Gen II) using the portal

Home > New > Storage account > Create storage account

Create storage account

Basics Networking Advanced Tags Review + create

Security

Secure transfer required Disabled Enabled

Azure Files

Large file shares Disabled Enabled

The current combination of storage account kind, performance, replication and location does not support large file shares.

Data protection

Blob soft delete Disabled Enabled

Data protection and hierarchical namespace cannot be enabled simultaneously.

Data Lake Storage Gen2

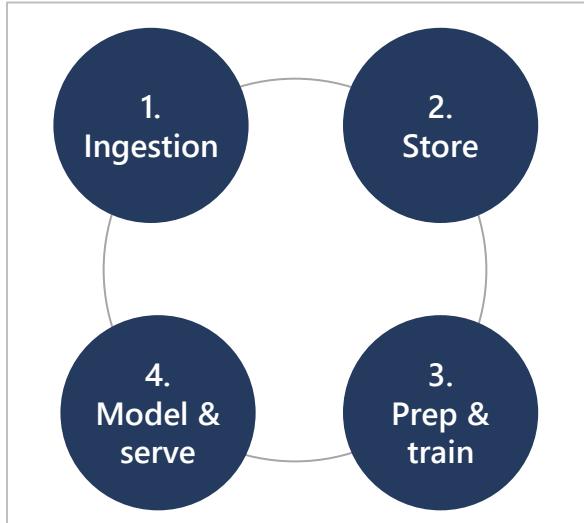
Hierarchical namespace Disabled Enabled

NFS v3 Disabled Enabled

Signup is currently required to utilize the NFS v3 feature on a per-subscription basis. [Signup for NFS v3](#)

46

Processing Big Data with Azure Data Lake Store



47

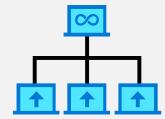
Big Data use cases

Let's examine three use cases for leveraging an Azure Data Lake Store

Modern data warehouse	This architecture sees Azure Data Lake Storage at the heart of the solution for a modern data warehouse. Using Azure Data Factory to ingest data into the Data Lake from a business application, and predictive models built in Azure Databricks, using Azure Synapse Analytics as a serving layer
Advanced analytics	In this solution, Azure Data factory is transferring terabytes of web logs from a web server to the Data Lake on an hourly basis. This data is provided as features to the predictive model in Azure Databricks, which is then trained and scored. The result of the model is then distributed globally using Azure Cosmos DB, that an application uses
Real time analytics	In this architecture, there are two ingestion streams. Azure Data Factory is used to ingest the summary files that are generated when the HGV engine is turned off. Apache Kafka provides the real-time ingestion engine for the telemetry data. Both data streams are stored in Data Lake store for use in the future

48

Lesson 04: Upload data into Azure Data Lake Store



49

Lesson objectives



Create an Azure Data Lake Gen2 Store using PowerShell



Upload data into the Data Lake Storage Gen2 using Azure Storage Explorer



Copy data from an Azure Data Lake Store Gen1 to an Azure Data Lake Store Gen2

50

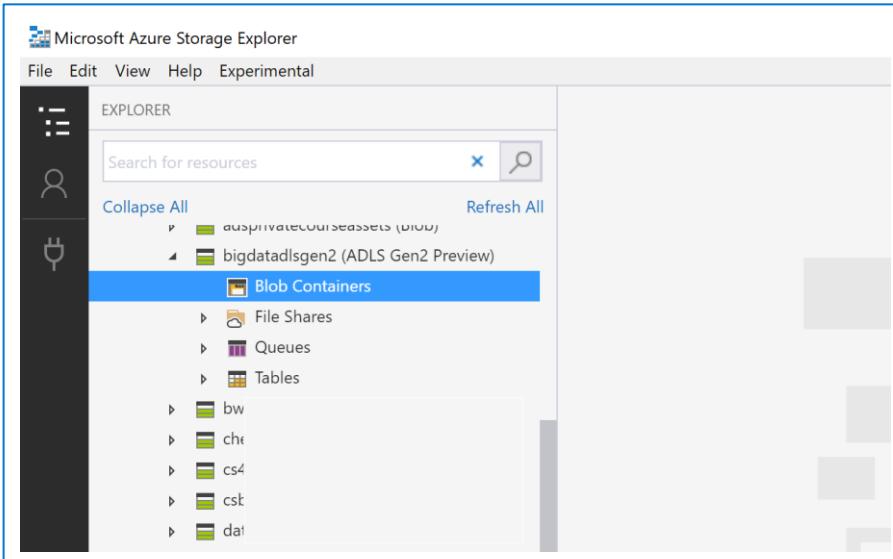
Create a Azure Data Lake Store (Gen II) using PowerShell

```
Windows PowerShell
Copyright (C) Microsoft Corporation. All rights reserved.

PS C:\Users> $location = "westus2"
>>
>> New-AzStorageAccount -ResourceGroupName $resourceGroup
>> -Name "storagequickstart"
>> -Location $location
>> -SkuName Standard_LRS
>> -Kind StorageV2
>> -EnableHierarchicalNamespace $True
```

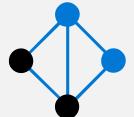
51

Uploading data with Azure Storage Explorer



52

Lab: Working with data storage



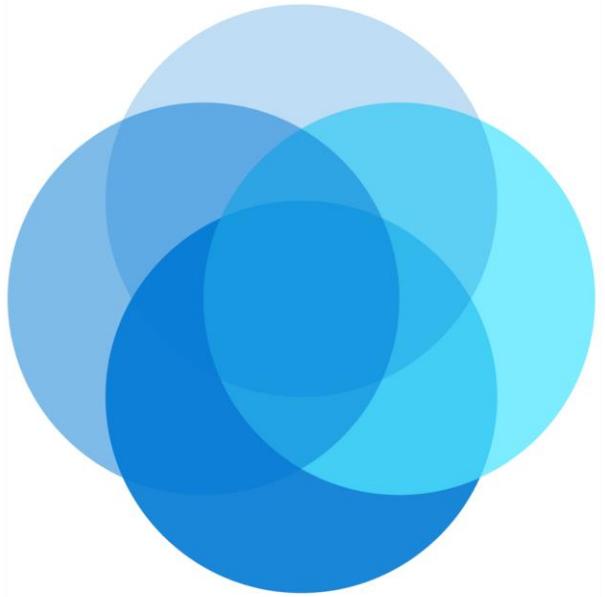
53

 Microsoft Azure

© Copyright Microsoft Corporation. All rights reserved.

54

Module 03: Enabling team based Data Science with Azure Databricks



55

Agenda



Lesson 01 – Describe Azure Databricks



Lesson 02 – Provision Azure Databricks and Workspaces



Lesson 03 – Read data using Azure Databricks



Lesson 04 – Perform transformations with Azure Databricks

56

Lesson 01: Describe Azure Databricks



57

Lesson objectives



What is Azure Databricks



What are Spark based analytics platform



How Azure Databricks integrates with enterprise security



How Azure Databricks integrates with other cloud services

58

What is Azure Databricks



Apache Spark-based analytics platform:

Simplifies the provisioning and collaboration of Apache Spark-based analytical solutions



Enterprise Security:

Utilizes the security capabilities of Azure



Integration with other Cloud Services:

Can integrate with a variety of Azure data platform services and Power BI

59

What is Apache Spark

Apache Spark emerged to provide a parallel processing framework that supports in-memory processing to boost the performance of big-data analytical applications on massive volumes of data

Interactive Data Analysis:

Used by business analysts or data engineers to analyze and prepare data

Streaming Analytics:

Ingest data from technologies such as Kafka and Flume to ingest data in real-time

Machine Learning:

Contains a number of libraries that enables a Data Scientist to perform Machine Learning

Why use Azure Databricks?

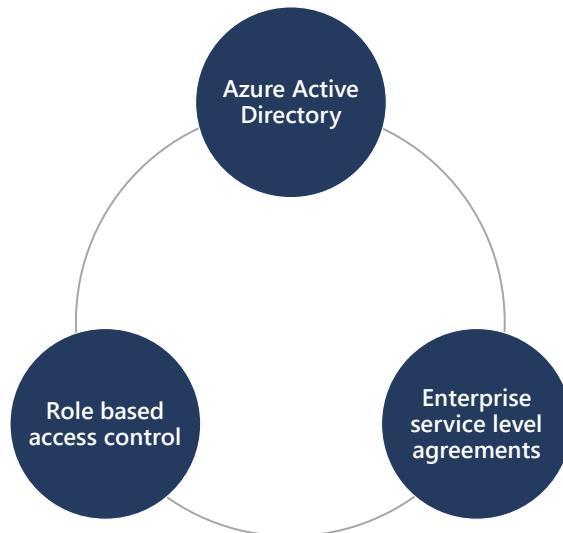
Azure Databricks is a wrapper around Apache Spark that simplifies the provisioning and configuration of a Spark cluster in a GUI interface

Azure Databricks components:

- Spark SQL and DataFrames
- Streaming
- Mlib
- GraphX
- Spark Core API

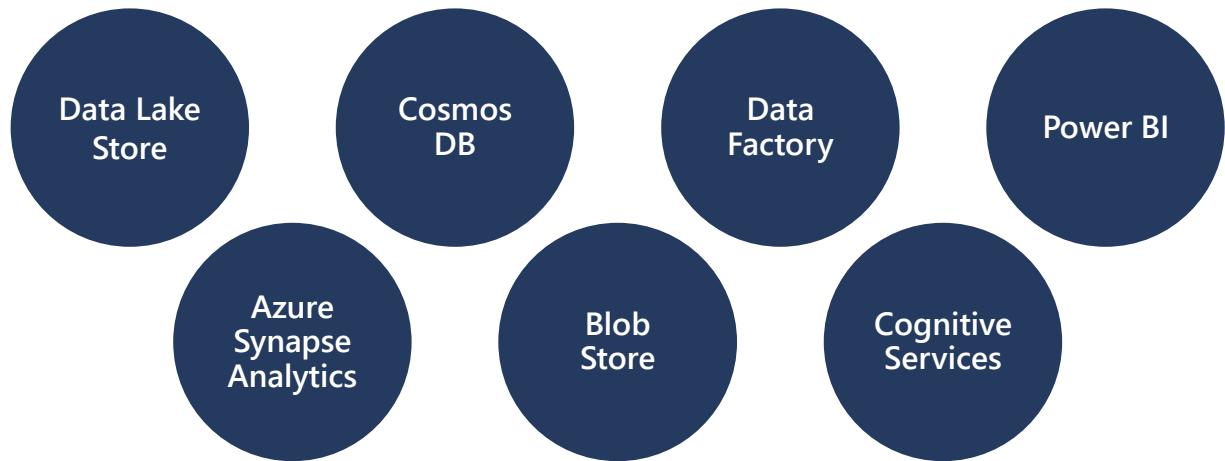
60

Enterprise security



61

Integration with cloud services



62

Lesson 02: Provision Azure Databricks and Workspaces



63

Lesson objectives



Create your own Azure Databricks workspace



Create a cluster and notebook in Azure Databricks

64

Create an Azure Databricks Workspace

[Home](#) > [New](#) > [Azure Databricks](#) > [Azure Databricks Service](#)

Azure Databricks Service

* Workspace name: ds-mslearn

* Subscription: (dropdown menu)

* Resource group: Create new (radio button selected) Use existing (radio button) cto_rg

* Location: West Europe

* Pricing Tier: Standard (Apache Spark, Secure with Azur...)

Deploy Azure Databricks workspace in your Virtual Network (preview)

Yes No

65

Create a Cluster and Notebook in Azure Databricks

Microsoft Azure PORTAL @microsoft.com

Azure Databricks

Explore the Quickstart Tutorial: Spin up a cluster, run queries on preloaded data, and display results in 5 minutes.

Import & Explore Data: Quickly import data, preview its schema, create a table, and query it in a notebook.

Create a Blank Notebook: Create a notebook to start querying, visualizing, and modeling your data.

Common Tasks

- New Notebook
- Upload Data
- Create Table
- New Cluster
- New Job
- New Mlflow Experiment New
- Import Library
- Read Documentation

Recents: Recent files appear here as you work.

Documentation

- Databricks Guide
- Python, R, Scala, SQL
- Importing Data

66

Lesson 03: Read data using Azure Databricks



67

Lesson objectives



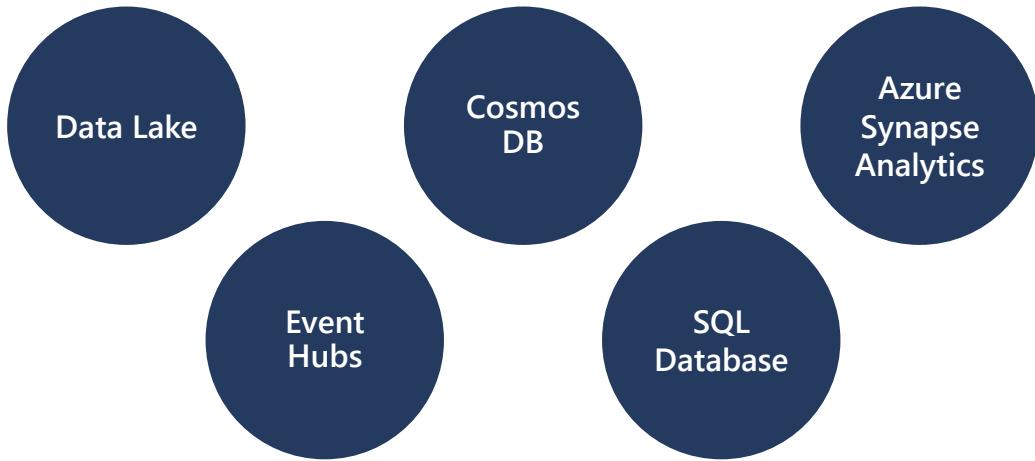
Use Azure Databricks to access data sources



Reading data in Azure Databricks

68

Use Azure Databricks to access data sources



69

Reading data in Azure Databricks

SQL	DataFrame (Python)
SELECT col_1 FROM myTable	df.select(col("col_1"))
DESCRIBE myTable	df.printSchema()
SELECT * FROM myTable WHERE col_1 > 0	df.filter(col("col_1") > 0)
..GROUP BY col_2	..groupBy(col("col_2"))
..ORDER BY col_2	..orderBy(col("col_2"))
..WHERE year(col_3) > 1990	..filter(year(col("col_3")) > 1990)
SELECT * FROM myTable LIMIT 10	df.limit(10)
display(myTable) (text format)	df.show()
display(myTable) (html format)	display(df)

70

Lesson 04: Perform transformations with Azure Databricks



71

Lesson objectives



Performing ETL to populate a data model



Perform basic transformations



Perform advanced transformations with user-defined functions

72

Performing ETL to populate a data model

The goal of transformation in Extract Transform Load (ETL) is to transform raw data to populate a data model

Extraction	Data validation	Transformation	Corrupt record handling	Loading data
Connect to many data stores: Postgres SQL Server Cassandra Cosmos DB CSV, Parquet Many more..	Validate that the data is what you expect	Applying structure and schema to your data to transform it into the desired format	Built-in functions of Databricks allow you to handle corrupt data such as missing and incomplete information	Highly effective design pattern involves loading structured data back to DBFS as a parquet file

73



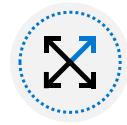
Normalizing values



Missing/null data



De-duplication



Pivoting data frames

74

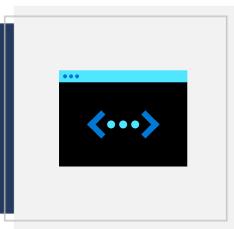
Advanced transformations

Advanced data transformation using custom and advanced user-defined functions, managing complex tables and loading data into multiple databases simultaneously

User-defined functions	This fulfils scenarios when you need to define logic specific to your use case and when you need to encapsulate that solution for reuse. UDFs provide custom, generalizable code that you can apply to ETL workloads when Spark's built-in functions won't suffice
Joins and lookup tables	A standard (or shuffle) join moves all the data on the cluster for each table to a given node on the cluster. This is an expensive operation. Broadcast joins remedy this situation when one DataFrame is sufficiently small enough to duplicate on each node of the cluster, avoiding the cost of shuffling a bigger DataFrame
Multiple databases	Loading transformed data to multiple target databases can be a time-consuming activity. Partitions and slots are options to get optimum performance from database connections. A partition refers to the distribution of data while a slot refers to the distribution of computation

75

Lab: Enabling team-based Data Science with Azure Databricks



76

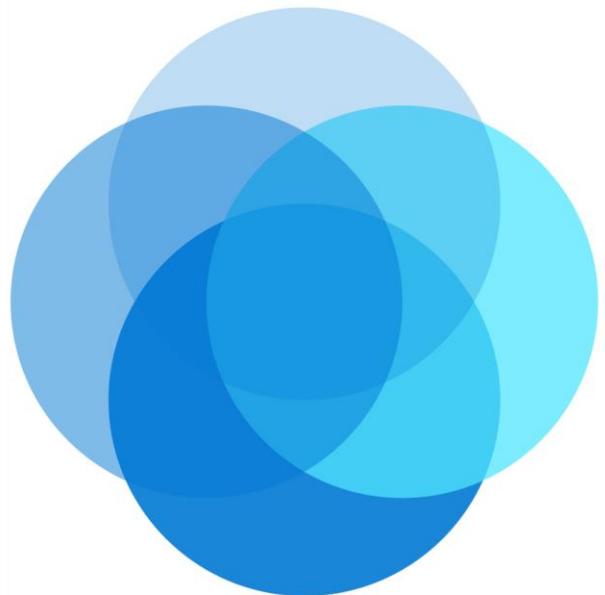


© Copyright Microsoft Corporation. All rights reserved.

77



Module 04: Building globally distributed databases with Cosmos DB



78

Agenda



Lesson 01: Create an Azure Cosmos DB database built to scale



Lesson 02: Insert and query data in your Azure Cosmos DB database



Lesson 03: Build a .NET Core app for Azure Cosmos DB in Visual Studio Code



Lesson 04: Distribute your data globally with Azure Cosmos DB

79

Lesson 01: Create an Azure Cosmos DB database built to scale



80

Lesson objectives



What is Cosmos DB



Create an Azure Cosmos DB account



What is a Request Unit



Choose a partition key



Create a database and container for NoSQL data in Azure Cosmos DB

81

What is Azure Cosmos DB



Scalability



Performance



Availability



Programming model

82

Create an Azure Cosmos DB account

83

What are Request Units

Throughput is important to ensure you can handle the volume of transactions you need

Database throughput	Database throughput is the number of reads and writes that your database can perform in a single second
What is a Request Unit	Azure Cosmos DB measures throughput using something called a request unit (RU). Request unit usage is measured per second, so the unit of measure is request units per second (RU/s). You must reserve the number of RU/s you want Azure Cosmos DB to provision in advance
Exceeding throughput limits	If you don't reserve enough request units, and you attempt to read or write more data than your provisioned throughput allows, your request will be rate-limited

84

Item size	Reads/second	Writes/second	Request units
1 KB	500	100	$(500 * 1) + (100 * 5) = 1,000 \text{ RU/s}$
1 KB	500	500	$(500 * 1) + (500 * 5) = 3,000 \text{ RU/s}$
4 KB	500	100	$(500 * 1.3) + (100 * 7) = 1,350 \text{ RU/s}$
4 KB	500	500	$(500 * 1.3) + (500 * 7) = 4,150 \text{ RU/s}$
64 KB	500	100	$(500 * 10) + (100 * 48) = 9,800 \text{ RU/s}$
64 KB	500	500	$(500 * 10) + (500 * 48) = 29,000 \text{ RU/s}$

85

Choosing a Partition Key

Why have a Partition Strategy?
 Having a partition strategy ensures that when your database needs to grow, it can do so easily and continue to perform efficient queries and transactions

What is a Partition Key?
 A partition key is the value by which azure organizes your data into logical divisions

86

Creating a Database and a Container in Cosmos DB

The screenshot shows the 'Add Container' dialog box. At the top, there's a message: 'Start at \$24/mo per database, multiple containers included' with a 'More details' link. Below that, the 'Database id' section has a radio button for 'Create new' (selected) and 'Use existing'. A text input field is present for 'Type a new database id'. The 'Provision database throughput' section contains a checked checkbox for 'Throughput (400 - 100,000 RU/s)' and radio buttons for 'Autopilot (preview)' (selected) and 'Manual'. A throughput value of '400' is entered. An estimated spend note follows: 'Estimated spend (USD): \$0.032 hourly / \$0.77 daily (1 region, 400RU/s, \$0.00008/RU)'. The 'Container id' section has a text input field with 'e.g., Container1'. The 'Partition key' section has a text input field with 'e.g., /address/zipCode'. A note below says 'My partition key is larger than 100 bytes'. The 'Unique keys' section has a plus sign and a 'Add unique key' link.

87

Lesson 02: Insert and Query Data in your Azure Cosmos DB Database



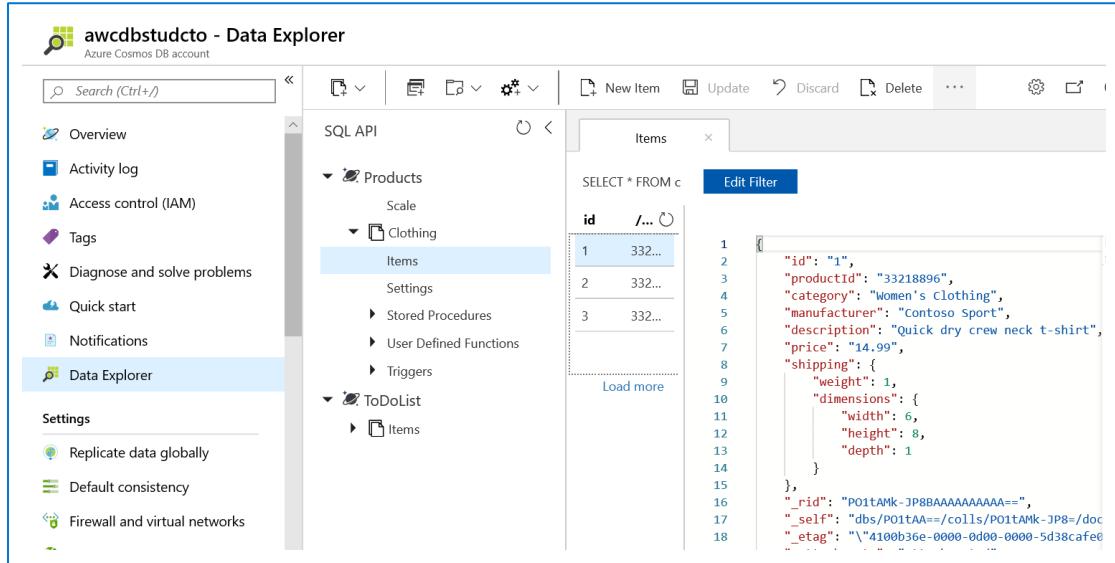
88

Lesson objectives

-  Create a product catalog document in the Data Explorer:
Add data
-  Perform Azure Cosmos DB queries:
Query types
Run queries
-  Running complex operations on your data
-  Working with graph data

89

Create a product catalog documents in the Data Explorer



The screenshot shows the Azure Cosmos DB Data Explorer interface for the account "awcdbstudcto". The left sidebar includes links for Overview, Activity log, Access control (IAM), Tags, Diagnose and solve problems, Quick start, Notifications, and Data Explorer (which is selected). Under Settings, there are options for Replicate data globally, Default consistency, and Firewall and virtual networks. The main area displays a SQL API query results pane titled "Items" with a table showing three items with IDs 1, 2, and 3. To the right, the raw JSON document for item ID 1 is shown:

```

1  {
2    "id": "1",
3    "productId": "33218896",
4    "category": "Women's Clothing",
5    "manufacturer": "Contoso Sport",
6    "description": "Quick dry crew neck t-shirt",
7    "price": "14.99",
8    "shipping": {
9      "weight": 1,
10     "dimensions": {
11       "width": 6,
12       "height": 8,
13       "depth": 1
14     }
15   },
16   "_rid": "P01tAMk-JP8BAAAAAAA=",
17   "_self": "dbs/P01tAa==/colls/P01tAMk-JP8=/doc",
18   "_etag": "\"4100b36e-0000-0d00-0000-5d38cafe0"
  
```

90

Perform Azure Cosmos DB Queries

SELECT Query Basics

```
SELECT <select_list>
[FROM <optional_from_specification>]
[WHERE <optional_filter_condition>]
[ORDER BY <optional_sort_specification>]
[JOIN <optional_join_specification>]
```

Examples

```
SELECT *
FROM Products p WHERE p.id ="1"
SELECT p.id, p.manufacturer, p.description
FROM Products p WHERE p.id ="1"
SELECT p.price, p.description, p.productId
FROM Products p ORDER BY p.price ASC
SELECT p.productId
FROM Products p JOIN p.shipping
```

91

Running complex operations on data

Multiple documents in your database frequently need to be updated at the same time. The way to perform these transactions in Azure Cosmos DB is by using stored procedures and user-defined functions (UDFs)

Stored procedures

Stored procedures perform complex transactions on documents and properties. Stored procedures are written in JavaScript and are stored in a collection on Azure Cosmos DB

User defined functions

User Defined Functions are used to extend the Azure Cosmos DB SQL query language grammar and implement custom business logic, such as calculations on properties and documents

92

Working with Graph Data

```
from gremlin_python.driver import client,
serializer
import sys, traceback

CLEANUP_GRAPH = "g.V().drop()"

INSERT_NATIONAL_PARK_VERTICES = [
    "g.addV('Park').property('id',",
    "'p1'),property('name',",
    "'Yosemite').property('Feature', 'El Capitan')",
    ", g.addV('Park').property('id',",
    "'p2').property('name', 'Joshua",
    "Tree').property('Feature', 'Yucca Brevifolia')",
    ", g.addV('State').property('id',",
    "'s1').property('name',",
    "'California').property('Location', 'USA')",
    ", g.addV('Ecosystem').property('id',",
    "'e1').property('name', 'Alpine')",
    ", g.addV('Ecosystem').property('id',",
    "'e2').property('name', 'Desert')",
    ", g.addV('Ecosystem').property('id',",
    "'e3').property('name', 'High Altitude')",
]

INSERT_NATIONAL_PARK_EDGES = [
    "g.V('p1').addE('is in').to(g.V('s1'))",
    ", g.V('p2').addE('is in').to(g.V('s1'))",
    ", g.V('p1').addE('has ecosystem",
    "of').to(g.V('e1'))",
    ", g.V('p2').addE('has ecosystem",
    "of').to(g.V('e2'))",
    ", g.V('p1').addE('has ecosystem",
    "of').to(g.V('e3'))",
    ", g.V('p2').addE('has ecosystem",
    "of').to(g.V('e3'))"
]
```

93

Lesson 03: Build a .NET Core App for Azure Cosmos DB in VS Code



94

Lesson objectives



Create an Azure Cosmos DB account, database, and container in Visual Studio Code using the Azure Cosmos DB extension



Create an application to store and query data in Azure Cosmos DB



Use the Terminal in Visual Studio Code to quickly create a console application



Add Azure Cosmos DB functionality with the help of the Azure Cosmos DB extension for Visual Studio Code

95

Creating Azure Cosmos DB in Visual Studio Code

```

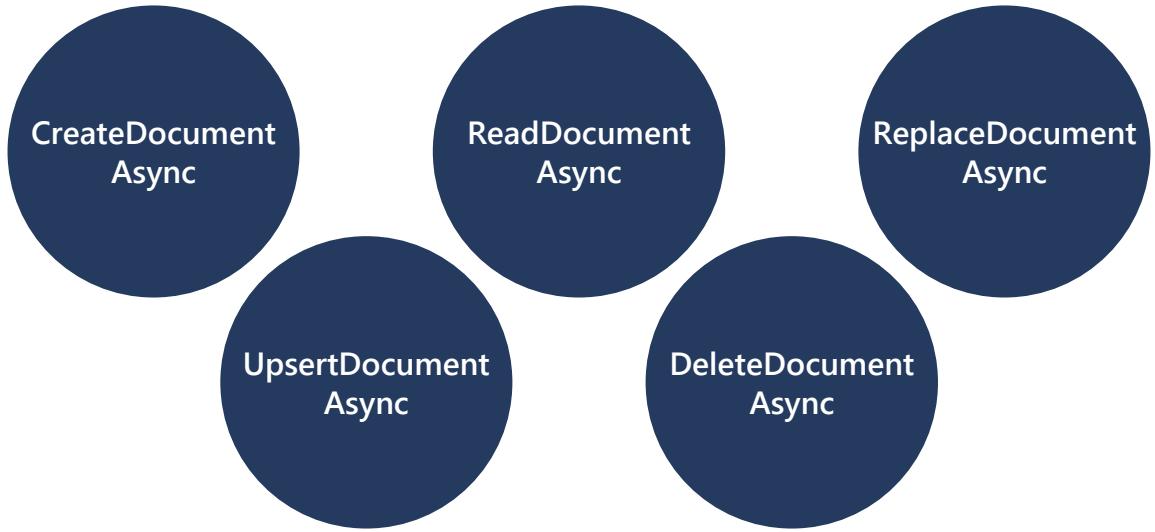
AZURE
COSMOS DB
adventbikes (MongoDB)
ctocdb (SQL)
Products
  Clothing
    Documents
      1
        Stored Procedures
      ToDoList
Attached Database Accounts

File Edit Selection View Go Debug Terminal Help
1-cosmos-document.json - Visual Studio Code
document.json
Press 'Enter' to confirm your input or 'Escape' to cancel
{
  "productId": "33218896",
  "category": "Women's Clothing",
  "manufacturer": "Contoso Sport",
  "description": "Quick dry crew neck t-shirt",
  "price": "14.99",
  "shipping": {
    "weight": 1,
    "dimensions": {
      "width": 6,
      "height": 8,
      "depth": 1
    }
  },
  "_rid": "QSl9ALDRxXUBAAAAAAA==",
  "_self": "dbs/QSl9AA==/colls/QSl9ALDRxXU=/docs/QSl9ALDRxXU=/1",
  "_etag": "\\"13000b6a-0000-0700-0000-5c9b619",
  "_attachments": "attachments/",
  "_ts": 1553686941
}

```

96

Working with documents programmatically



97

Querying document programmatically

```

{
  // Set some common query options
  FeedOptions queryOptions = new FeedOptions { MaxItemCount = -1, EnableCrossPartitionQuery = true };

  // Here we find nelapin via their LastName
  IQueryables<User> userQuery = this.client.CreateDocumentQuery<User>(
    UriFactory.CreateDocumentCollectionUri(databaseName, collectionName), queryOptions)
    .Where(u => u.LastName == "Pindakova");

  // The query is executed synchronously here, but can also be executed asynchronously via the IDocumentQuery<T> interface
  Console.WriteLine("Running LINQ query...");
  foreach (User in userQuery)
  {
    Console.WriteLine("\tRead {0}", user);
  }

  // Now execute the same query via direct SQL
  IQueryables<User> userQueryInSql = this.client.CreateDocumentQuery<User>(
    UriFactory.CreateDocumentCollectionUri(databaseName, collectionName),
    "SELECT * FROM User WHERE User.lastName = 'Pindakova'", queryOptions );

  Console.WriteLine("Running direct SQL query...");
  foreach (User in userQueryInSql)
  {
    Console.WriteLine("\tRead {0}", user);
  }

  Console.WriteLine("Press any key to continue ...");
  Console.ReadKey();
}
  
```

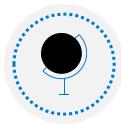
98

Lesson 04: Distribute your data globally with Azure Cosmos DB



99

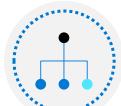
Lesson objectives



Learn about the benefits of writing and replicating data to multiple regions around the world



Cosmos DB multi-master replication



Cosmos DB failover management



Change the consistency setting for your database

100

Benefits of writing and replicating data to multiple regions

Home > Resource groups > cto_rg > ctocdb > Replicate data globally

Replicate data globally

ctocdb

Save Discard Manual Failover Automatic Failover

Click on a location to add or remove regions from your Azure Cosmos DB account.

* Each region is billable based on the throughput and storage for the account. [Learn more](#)

Configure regions

Configure the regions available for reads and writes. + Add region

REGIONS	READS ENABLED	WRITES ENABLED	
West US	✓	✓	trash
UK South	✓	✓	trash
Japan West	✓	✓	trash
South Africa North	✓	✓	trash

101

Cosmos DB multi-master replication



102

Cosmos DB failover management

Automated fail-over is a feature that comes into play when there's a disaster or other event that takes one of your read or write regions offline, and it redirects requests from the offline region to the next most prioritized region

Read region outage

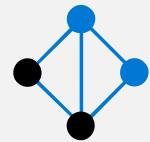
Azure Cosmos DB accounts with a read region in one of the affected regions are automatically disconnected from their write region and marked offline

Write region outage

If the affected region is the current write region and automatic fail-over is enabled, then the region is automatically marked as offline. Then, an alternative region is promoted as the write region

103

Lab: Building globally distributed databases with Cosmos DB

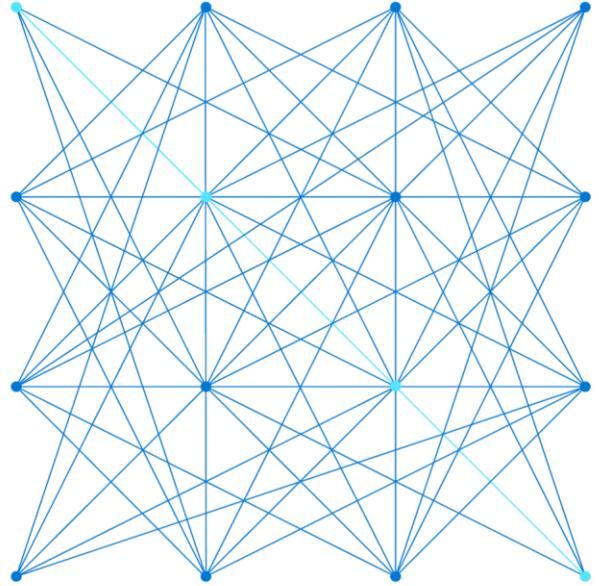


104

© Copyright Microsoft Corporation. All rights reserved.

105

Module 05: Working with relational data stores in the cloud

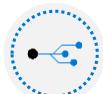


106

Agenda



Lesson 01: Work with Azure SQL database



Lesson 02: Work with Azure Synapse Analytics



Lesson 03: Provision and query data in Azure Synapse Analytics



Lesson 04: Import data into Azure Synapse Analytics using PolyBase

107

Lesson 01: Azure SQL database



108

Lesson objectives



Why Azure SQL Database is a good choice
for running your relational database



What configuration and pricing options
are available for your Azure SQL database



How to create an Azure SQL database from
the portal



How to use Azure Cloud Shell to connect to
your Azure SQL database, add a table, and work
with data

109

Why Azure SQL Database is a good choice



Convenience



Cost



Scale



Security

110

Azure SQL Database configuration options

When you create your first Azure SQL database, you also create an *Azure SQL logical server*. Think of a logical server as an administrative container for your databases

DTUs	vCores	SQL elastic pools	SQL managed instances
DTU stands for Database Transaction Unit and is a combined measure of compute, storage, and IO resources. Think of the DTU model as a simple, preconfigured purchase option	vCore gives you greater control over what compute and storage resources you create and pay for. vCore model enables you to configure resources independently	SQL elastic pools relate to eDTUs. They enable you to buy a set of compute and storage resources that are shared among all the databases in the pool. Each database can use the resources they need	The SQL managed instance creates a database with near 100% compatibility with the latest SQL Server on-premises Enterprise Edition database engine, useful for SQL Server customers who would like to migrate on-premises servers instance in a "lift and shift" manner

111

Create an Azure SQL database

Home > New > SQL Database > Create SQL Database

Create SQL Database

Basics Networking Additional settings Tags Review + create

Create a SQL database with your preferred configurations. Complete the Basics tab then go to Review + Create to provision with smart defaults, or visit each tab to customize. [Learn more](#)

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription * chtestao

Resource group * Select existing... Create new

Database details

Enter required settings for this database, including picking a logical server and configuring the compute and storage resources

Database name * Enter database name

Server * Create new

The value must not be empty.

Want to use SQL elastic pool? * Yes No

Compute + storage * Configure database

Dashboard > New > Create SQL Database

1 Create SQL Database

2 Basics Additional settings Tags Review + create

Customize additional configuration parameters including collation & sample data.

Data source

Start with a blank database, restore from a backup or select sample data to populate your new database.

* Use existing data

* Backup

3 None Backup Sample

4 Select a backup

myserver (West Europe)

- database1 (2019-09-16 12:05:30 UTC)
- database2 (2019-09-16 12:06:45 UTC)
- database3 (2019-09-16 12:07:51 UTC)
- database4 (2019-09-16 12:08:38 UTC)
- database5 (2019-09-16 12:09:23 UTC)
- database6 (2019-09-16 12:10:41 UTC)
- database7 (2019-09-16 12:11:38 UTC)

Database Collation

Database collation defines the rules that sort and compare data, and cannot be changed after database creation. The default database collation is SQL_Latin1_General_CI_AS. [Learn more](#)

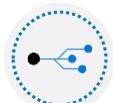
112

Lesson 02: Azure Synapse Analytics



113

Lesson objectives



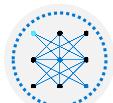
Explain Azure Synapse Analytics



Explain Azure Synapse Analytics features



Types of solution workloads



Explain massively parallel processing concepts



Compare table geometries

114

Azure Synapse Analytics

What is Azure Synapse Analytics?

A unified environment by combining the enterprise data warehouse of SQL, the Big Data analytics capabilities of Spark, and data integration technologies to ease the movement of data between both, and from external data sources

Data warehouse capabilities

SQL Analytics:

A centralized data warehouse store that provides a relational analytics and decision support services across the whole enterprise

SQL Pools:

CPU, memory, and IO are bundled into units of compute scale called SQL, determined by Data Warehousing Units (DWU)

Future features:

Will include a Spark engine, a data integration and Azure Synapse Analytics Studio

115

Azure Synapse Analytics features

Workload management

This capability is used to prioritize the query workloads that take place on the server using Workload Management. This involves three components:

- Workload Groups
- Workload Classification
- Workload Importance

Result-set cache

Result-set caching can be used to improve the performance of the queries that retrieve these results. When result-set caching is enabled, the results of the query are cached in the SQL pool storage

Materialized views

A materialized view pre-computes, stores, and maintains its data like a table. They are automatically updated when data in underlying tables are changed

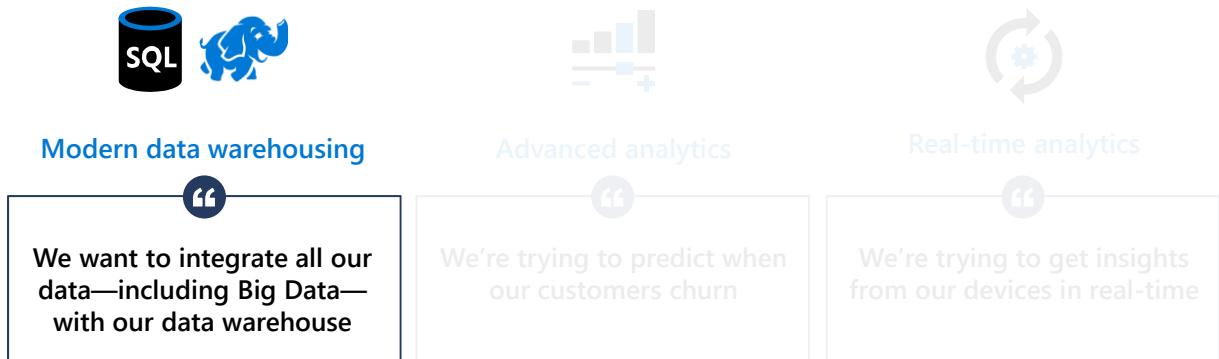
SSDT CI/CD support

Database project support in SQL Server Data Tools (SSDT) allows teams of developers to collaborate over a version-controlled Azure Synapse Analytics, and track, deploy and test schema changes

116

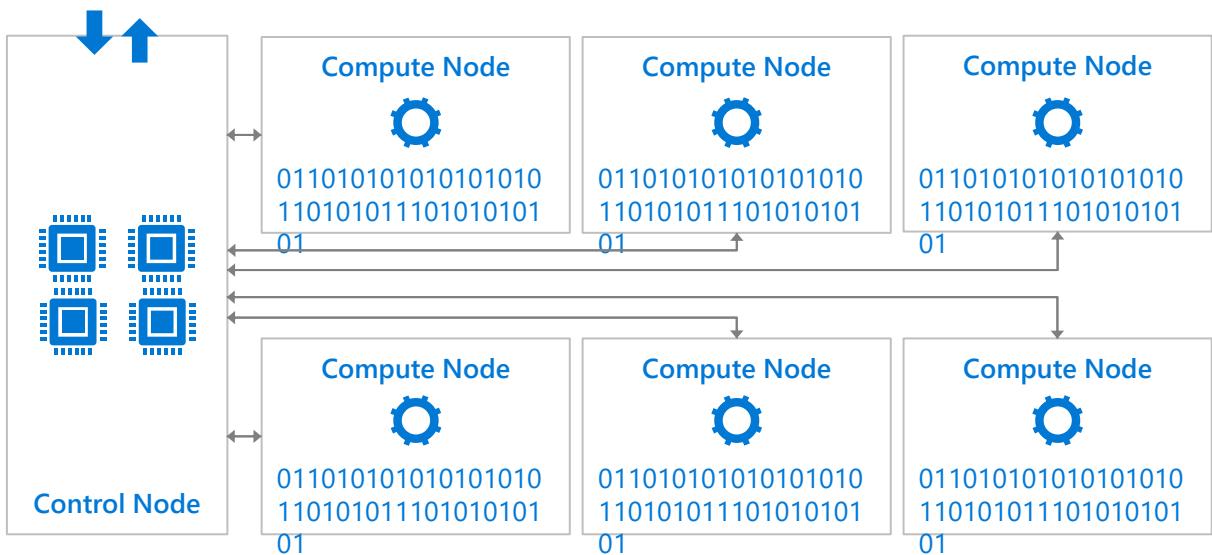
Types of solution workloads

The modern data warehouse extends the scope of the data warehouse to serve Big Data that's prepared with techniques beyond relational ETL



117

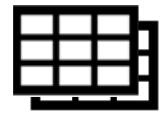
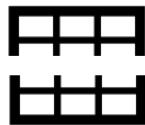
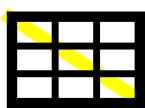
Massively Parallel Processing (MPP) concepts



118

Table geometries

Table distribution



Round Robin Tables

Hash Distributed Tables

Replicated Tables

119

Lesson 03: Creating and querying an Azure Synapse Analytics



120



Lesson objectives



Create an Azure Synapse Analytics sample database



Query the sample database with the SELECT statement and its clauses



Use the queries in different client applications such as SQL Server Management Studio, and PowerBI

121

Create an Azure Synapse Analytics

The screenshot shows the Azure portal's 'Create a SQL data warehouse' wizard. The top navigation bar includes 'Home', 'New', 'Azure Synapse Analytics (formerly SQL DW)', and 'SQL Data Warehouse'. The main heading is 'Welcome to Azure Synapse Analytics (formerly known as Azure SQL Data Warehouse)'. Below it, there are tabs for 'Basics', 'Additional settings', 'Tags', and 'Review + create'. A note says 'Create a SQL data warehouse with your preferred configurations. Complete the Basics tab then go to Review + Create to provision with smart defaults, or visit each tab to customize.' Under 'Project details', there are fields for 'Subscription' (set to 'chtestao') and 'Resource group' (with options 'Select existing...' and 'Create new'). The 'Data warehouse details' section requires entering a 'Data warehouse name' and selecting a 'Server'. A validation error message 'The value must not be empty.' is shown next to the server selection field. The 'Performance level' section has a note 'Please select a server first.' A large blue border surrounds the entire form area.

122

Perform Azure Synapse Analytics queries

SELECT Query Basics

```
SELECT <select_list>
[FROM <optional_from_specification>]
[WHERE <optional_filter_condition>]
[ORDER BY <optional_sort_specification>]
[JOIN <optional_join_specification>]
```

Examples

```
SELECT *
FROM Products p WHERE p.id ="1"
```

```
SELECT p.id, p.manufacturer, p.description
FROM Products p WHERE p.id ="1"
```

```
SELECT p.price, p.description, p.productId
FROM Products p ORDER BY p.price ASC
```

```
SELECT p.productId
FROM Products p JOIN p.shipping
```

123

Perform Azure Synapse Analytics queries

Create Table as Select (CTAS)

Used in parallel data loads

CREATE TABLE

```
[ database_name . [ schema_name ] . | schema_name. ] table_name
    [ ( { column_name } [ ,...n ] ) ]
WITH ( DISTRIBUTION =
    { HASH( distribution_column_name )
        [ , <CTAS_table_option> [ ,...n ] ]
    }
    AS <select_statement> [;]
```

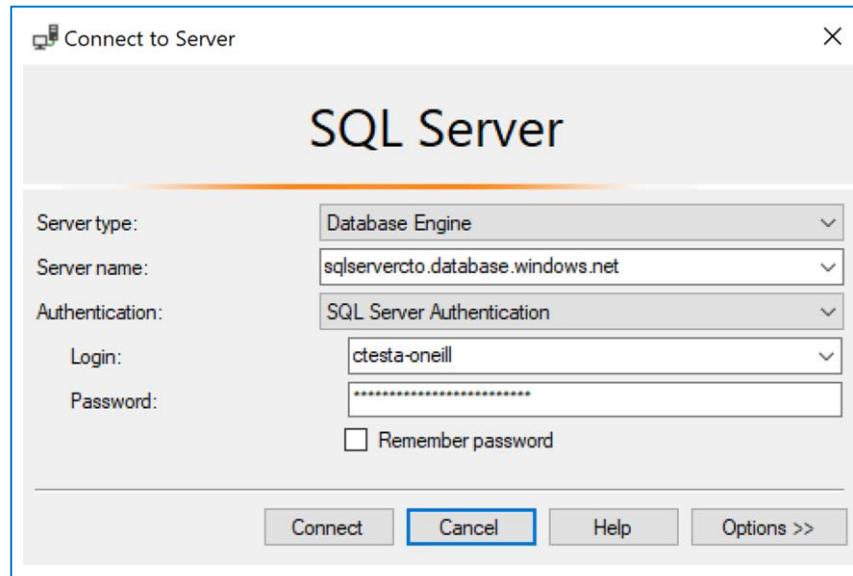
REPLICATE | ROUND_ROBIN }

Example

```
CREATE TABLE FactInternetSales_Copy
WITH
(DISTRIBUTION = HASH(SalesOrderNumber))
AS SELECT * FROM FactInternetSales
```

124

Querying with different client applications



125

Lesson 04: Using PolyBase to Load Data in Azure Synapse Analytics



126

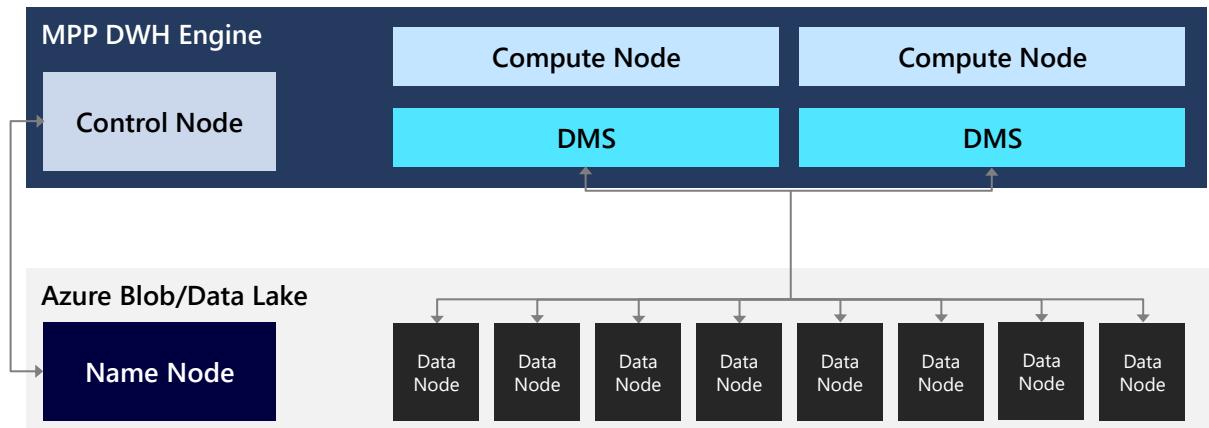
Lesson objectives

-  Explore how PolyBase works
-  Upload text data to Azure Blob store
-  Collect the security keys for Azure Blob store
-  Create an Azure Synapse Analytics
-  Import data from Blob Storage to the Data Warehouse

127

How PolyBase works

The MPP engine's integration method with PolyBase



128

Upload text data to Azure Blob store

The screenshot shows the 'Create storage account' wizard on the 'Basics' step. It includes fields for Subscription (selected), Resource group (selected), Storage account name (toazureblob), Location (West Europe), Performance (Standard selected), Account kind (StorageV2 selected), Replication (Read-access geo-redundant storage (RA-GRS) selected), and Access tier (Hot selected). Navigation buttons at the bottom are 'Review + create' (highlighted in blue), 'Previous', and 'Next : Advanced >'.

129

Collect the storage keys

The screenshot shows the 'Access keys' page for the 'toazureblob' storage account. It displays two sets of access keys: 'key1' and 'key2'. Each key has a 'Key' field containing a long, encoded string (e.g., eU7...Cg== for key1) and a 'Connection string' field. The connection strings begin with 'DefaultEndpointsProtocol=https;AccountName=toazureblob;AccountKey=' followed by the respective key's value. The 'key1' section also includes a note about regenerating keys and updating applications.

130

Create an Azure Synapse Analytics

Home > New > Azure Synapse Analytics (formerly SQL DW) > SQL Data Warehouse

SQL Data Warehouse

Welcome to Azure Synapse Analytics (formerly known as Azure SQL Data Warehouse). [Learn more](#)

Basics * [Additional settings](#) * [Tags](#) [Review + create](#)

Create a SQL data warehouse with your preferred configurations. Complete the Basics tab then go to Review + Create to provision with smart defaults, or visit each tab to customize. [Learn more](#)

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription * [chtestao](#)

Resource group * [Select existing...](#) [Create new](#)

Data warehouse details

Enter required settings for this data warehouse, including picking a logical server and configuring the performance level.

Data warehouse name *

Server * [Select a server](#) [Create new](#)

The value must not be empty.

Performance level * [Select performance level](#)

Please select a server first.

131

Lab: Working with relational data stores in the cloud



132

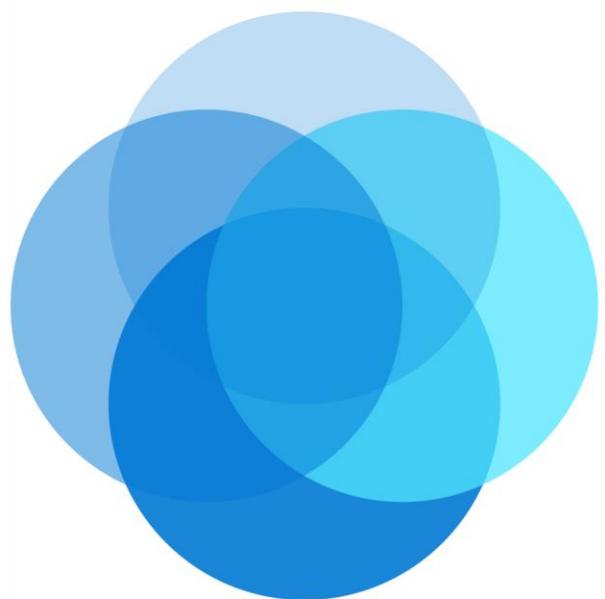


© Copyright Microsoft Corporation. All rights reserved.

133



Module 06: Performing real-time analytics with Stream Analytics



134

Agenda



Lesson 01 – Data streams and event processing



Lesson 02 – Data ingestion with Event Hubs



Lesson 03 – Processing data with Stream Analytics Jobs

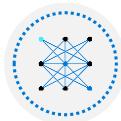
135

Lesson 01: Data streams and event processing

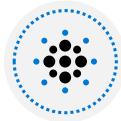


136

Lesson objectives



Explain data streams



Explain event processing



Learn about processing events with Azure Stream Analytics

137

What are data streams

Data streams:

In the context of analytics, data streams are event data generated by sensors or other sources that can be analyzed by another technology

Data stream processing approach:

There are two approaches. Reference data is streaming data that can be collected over time and persisted in storage as static data. In contrast, streaming data have relatively low storage requirements. And run computations in sliding windows

Data streams are used to:

Analyze data:
Continuously analyze data to detect issues and understand or respond to them

Understand systems:
Understand component or system behavior under various conditions to fuel further enhancements of said system

Trigger actions:
Trigger specific actions when certain thresholds are identified

138

Event processing

The process of consuming data streams, analyzing them, and deriving actionable insights out of them is called Event Processing and has three distinct components:

Event producer	Examples include sensors or processes that generate data continuously such as a heart rate monitor or a highway toll lane sensor
Event processor	An engine to consume event data streams and deriving insights from them. Depending on the problem space, event processors either process one incoming event at a time (such as a heart rate monitor) or process multiple events at a time (such as a highway toll lane sensor)
Event consumer	An application which consumes the data and takes specific action based on the insights. Examples of event consumers include alert generation, dashboards, or even sending data to another event processing engine

139

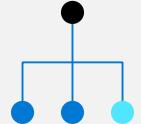
Processing events with Azure Stream Analytics

Microsoft Azure Stream Analytics is an event processing engine. It enables the consumption and analysis of high volumes of streaming data in real time

Source	Ingestion	Analytical engine	Destination
Sensors	Event Hubs	Stream Analytics Query Language	Azure Data Lake
Systems	IoT Hubs	.NET SDK	Cosmos DB
Applications	Azure Blob Store		SQL Database Blob Store Power BI

140

Lesson 02: Data ingestion with Event Hubs



141

Lesson objectives



Describe Azure Event Hubs



Create an Event Hub



Evaluate the performance of an Event Hub



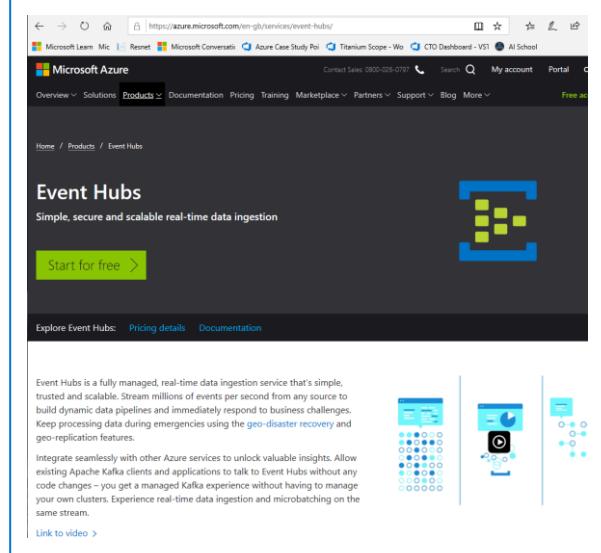
Configure applications to use an Event Hub

142

Azure Event Hubs



Azure Event Hubs is a highly scalable publish-subscribe service that can ingest millions of events per second and stream them into multiple applications



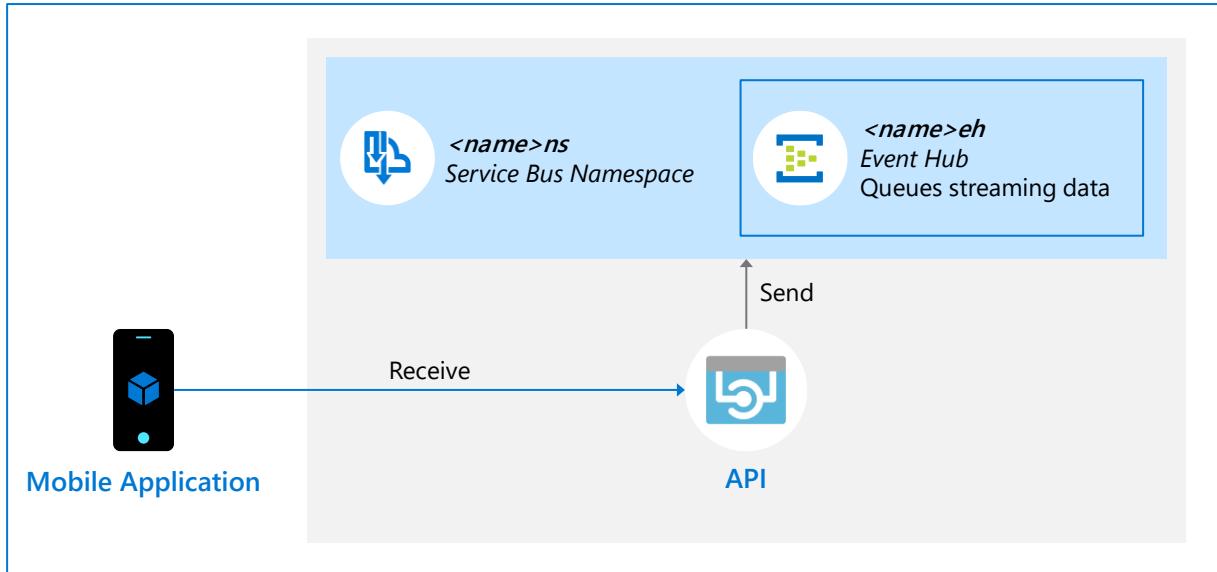
143

Create an Event Hub

Create an event hub namespace	Create an event hub
<ol style="list-style-type: none"> In the Azure portal, select NEW, type Event Hubs, and then select Event Hubs from the resulting search. Then select Create Provide a name for the event hub, and then create a resource group. Specify xx-name-eh and xx-name-rg respectively, XX- represent your initials to ensure uniqueness of the Event Hub name and Resource Group name Click the checkbox to Pin to the dashboard, then select the Create button 	<ol style="list-style-type: none"> After the deployment is complete, click the xx-name-eh event hub on the dashboard Then, under Entities, select Event Hubs To create the event hub, select the + Event Hub button. Provide the name socialstudy-eh, and then select Create To grant access to the event hub, we need to create a shared access policy. Select the socialstudy-eh event hub when it appears, and then, under Settings, select Shared access policies Under Shared access policies, create a policy with MANAGE permissions by selecting + Add. Give the policy the name of xx-name-eh-sap, check MANAGE, and then select Create Select your new policy after it has been created, and then select the copy button for the CONNECTION STRING – PRIMARY KEY entity Paste the CONNECTION STRING – PRIMARY KEY entity into Notepad, this is needed later in the exercise Leave all windows open

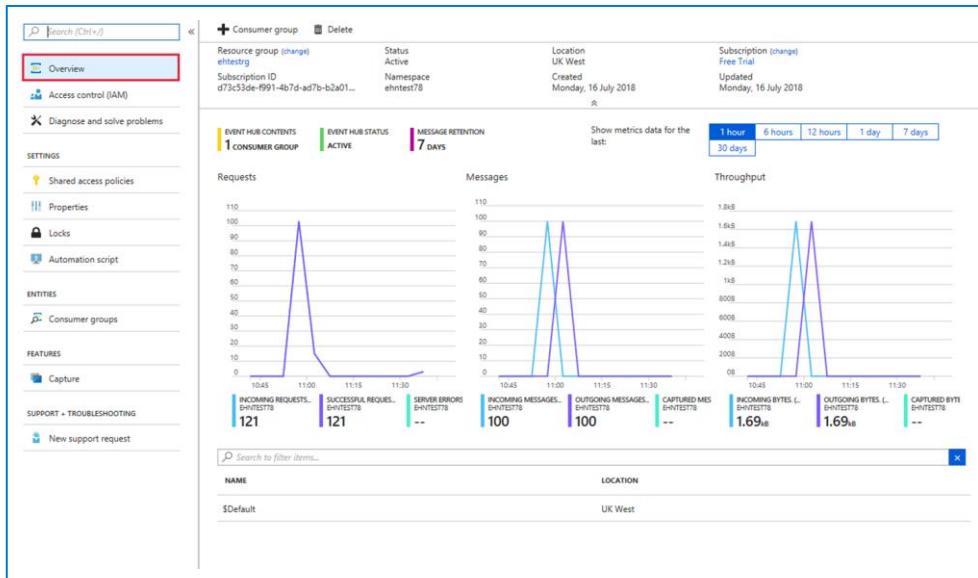
144

Configure applications to use Event Hubs



145

Evaluating the performance of Event Hubs



146

Lesson 03: Processing data with Stream Analytics Jobs



147

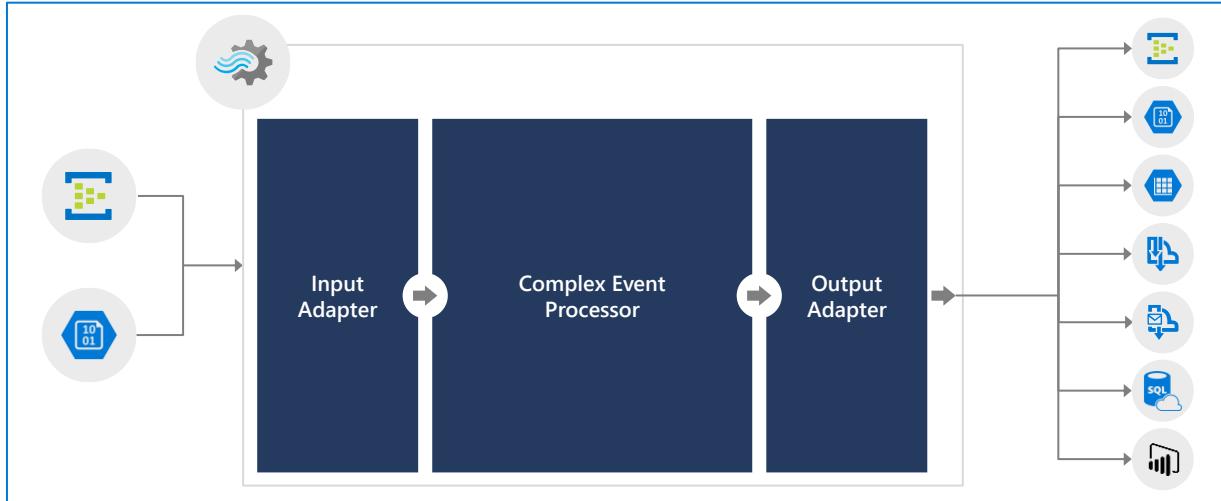
Lesson objectives

- Explore the Streaming Analytics workflow
- Create a Stream Analytics Job
- Configure a Stream Analytics job input
- Configure a Stream Analytics job output
- Write a transformation query
- Start a Stream Analytics job

148

Azure Stream Analytics workflow

Complex event processing of Stream Data in Azure



149

Create Stream Analytics service

Job name	Subscription	Resource group	Location
cto-as-a-job1		cto_rg	West Europe
<div style="border: 1px solid #ccc; padding: 10px;"> <p>Home > New > Stream Analytics job > New Stream Analytics job</p> <p>New Stream Analytics job</p> <p>* Job name: cto-as-a-job1</p> <p>* Subscription: (dropdown)</p> <p>* Resource group: cto_rg</p> <p>Create new</p> <p>* Location: West Europe</p> <p>Hosting environment: Cloud</p> <p>Streaming units (1 to 120): 6</p> </div>			

150

Create a Stream Analytics Job input

Event Hub
New input

* Input alias: cto-asa-input01

Provide Event Hub settings manually
 Select Event Hub from your subscriptions

Subscription: LearnAI Training Subscription

* Event Hub namespace: cto-eh-ns

* Event Hub name: cto-name-eh

* Event Hub policy name: RootManageSharedAccessKey

Event Hub policy key: *****

Event Hub consumer group:

* Event serialization format: JSON

Encoding: UTF-8

Event compression type: None

151

Create a Stream Analytics Job output

Home > Resource groups > cto_rg > cto-asa-job1 > Outputs

Outputs

+ Add

Event Hub	SINK
SQL Database	
Blob storage	
Table storage	
Service Bus topic	
Service Bus queue	
Cosmos DB	
Power BI	
Data Lake Storage Gen1	

Blob storage
New output

* Output alias: cto-asa-output01

Provide Blob storage settings manually
 Select Blob storage from your subscriptions

Subscription: LearnAI Training Subscription

* Storage account: ctoazureblob

* Storage account key: *****

* Container: socialmedia

Path pattern:

Date format: YYYY/MM/DD

Time format: HH

* Event serialization format: JSON

Encoding: UTF-8

152

Write a transformation query

The screenshot shows the Azure Stream Analytics job overview page for 'cto-asa-job1'. The 'Overview' section displays basic information: Resource group (change) : cto_rg, Status : Created, Location : West Europe, Subscription (change) : LearnAI Training Subscription, and Subscription ID : 5be49961-ea44-42ec-8021-b728be90d58c. On the right, there's a 'Query' pane containing the following T-SQL code:

```

1 SELECT
2 *
3 INTO
4 [cto-asa-output01]
5 FROM
6 [cto-asa-input01]

```

153

Start a Stream Analytics Job

The screenshot shows the Azure Stream Analytics job overview page for 'cto-asa-job1'. The 'Overview' section displays basic information: Resource group (change) : cto_rg, Status : Created, Location : West Europe, Subscription (change) : LearnAI Training Subscription, and Subscription ID : 5be49961-ea44-42ec-8021-b728be90d58c. The 'Start' button in the top navigation bar is highlighted with a red box. On the right, there's a 'Query' pane containing the same T-SQL code as the previous screenshot.

154

Lab: Performing real-time analytics with Stream Analytics



155

 Microsoft Azure

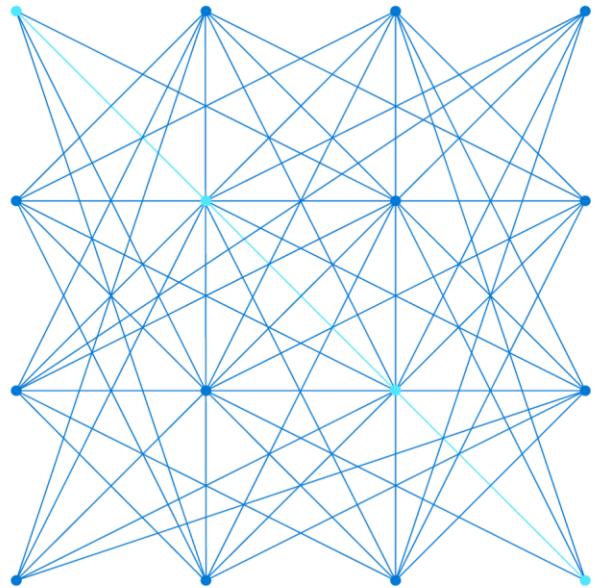
© Copyright Microsoft Corporation. All rights reserved.

156



Module 07: Orchestrating data movement with Azure Data Factory

Start : 09.15



157

Agenda

-  Lesson 01 – Introduction to Azure Data Factory
 -  Lesson 02 – Understand Azure Data Factory components
 -  Lesson 03 – Integrate Azure Data Factory with Databricks

158

Lesson 01: Introduction to Azure Data Factory



159

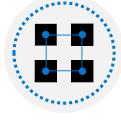
Lesson objectives



What is Azure Data Factory



The Data Factory process



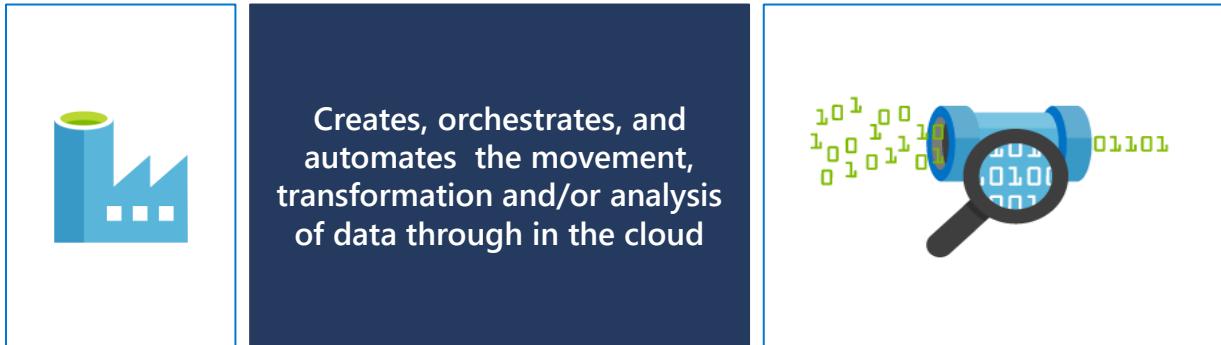
Azure Data Factory components



Azure Data Factory security

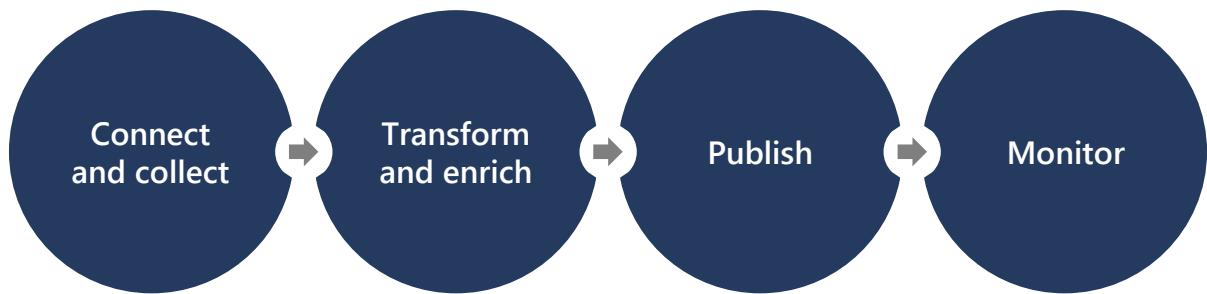
160

What is Azure Data Factory



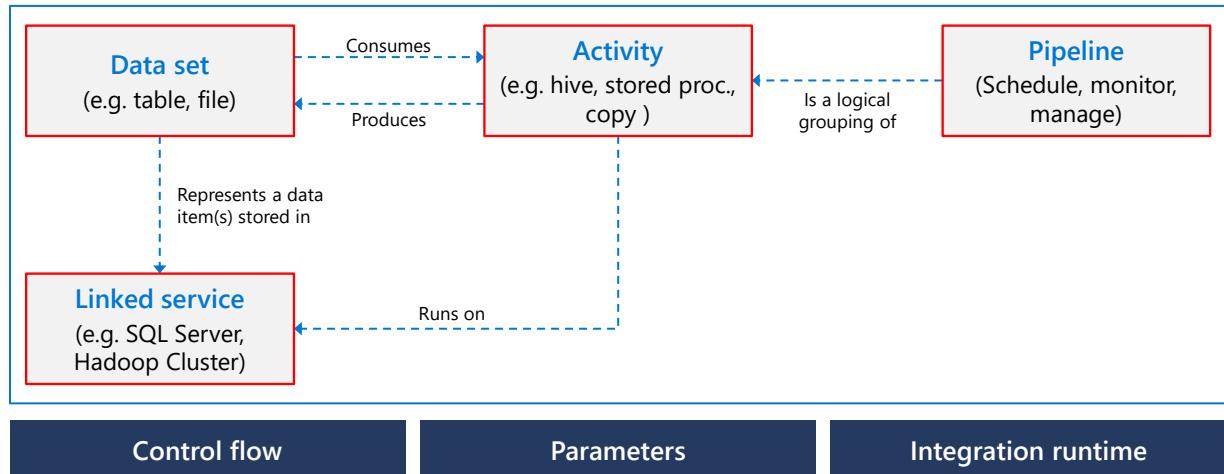
161

The Data Factory process



162

Azure Data Factory components



163

Azure Data Factory security

Data factory contributor role				
1	2	3	4	5
Create, edit, and delete data factories and child resources including datasets, linked services, pipelines, triggers, and integration runtimes	Deploy Resource Manager templates. Resource Manager deployment is the deployment method used by Data Factory in the Azure portal	Manage App Insights alerts for a data factory	At the resource group level or above, lets users deploy Resource Manager template	Create support tickets

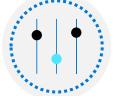
164

Lesson 02: Azure Data Factory components



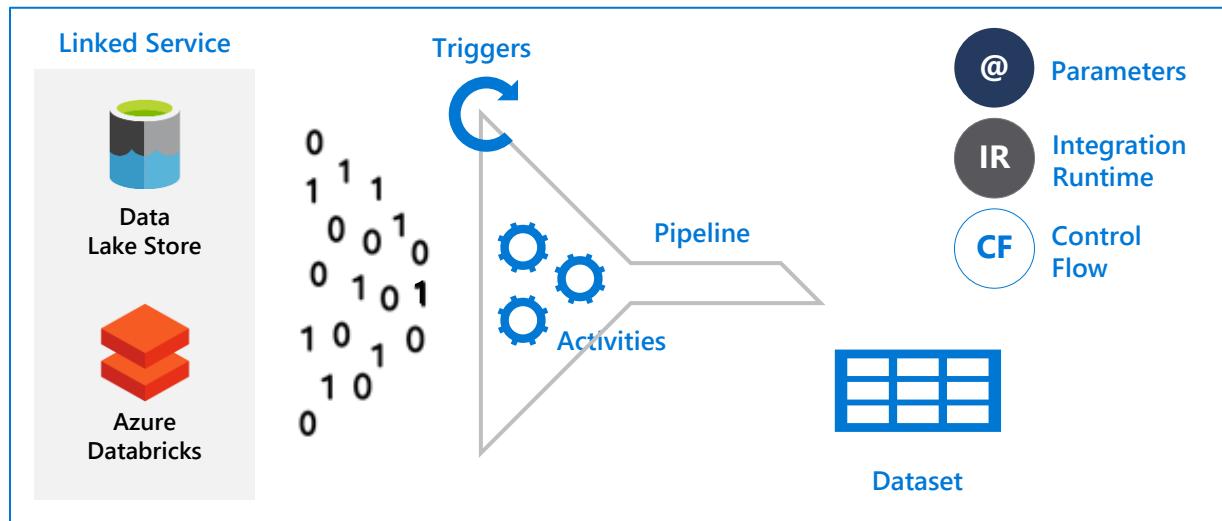
165

Lesson objectives

-  Linked services
-  Datasets
-  Data Factory activities
-  Pipelines
-  Pipeline example

166

Azure Data Factory components



167

Data factory activities

Activities within Azure Data Factory defines the actions that will be performed on the data and there are three categories including:

Data movement activities	Data movement activities simply move data from one data store to another. A common example of this is in using the Copy Activity
Data transformation activities	Data transformation activities use compute resource to change or enhance data through transformation, or it can call a compute resource to perform an analysis of the data
Control Activities	Control flow orchestrate pipeline activities that includes chaining activities in a sequence, branching, defining parameters at the pipeline level, and passing arguments while invoking the pipeline on-demand or from a trigger

168

Pipelines

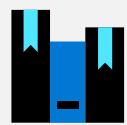
Pipeline is a grouping of logically related activities

Pipeline can be scheduled so the activities within it get executed

Pipeline can be managed and monitored

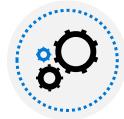
169

Lesson 03: Ingesting and transforming data



170

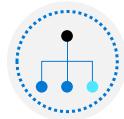
Lesson objectives



How to setup Azure Data Factory



Ingest data using the Copy Activity



Transforming data with the Mapping Data Flow

171

Create Azure Data Factory

Home > New > Data Factory > New data factory

New data factory

Name *

Version V2 V3

Subscription * chtestao

Resource Group *

Location * South Central US West Europe West US West US 2 East US East US 2 North Europe Central US Central US 2

Enable GIT

GIT URL *

Repo name *

Branch Name *

Root folder *

172

Ingesting data with the copy activity



Reads data from a source data store

Performs serialization/deserialization, compression/decompression, column mapping, and so on. It performs these operations based on the configuration of the input dataset, output dataset, and Copy activity

Writes data to the sink/destination data store

173

Transforming data with the Mapping Data Flow

Code free data transformation at scale

Perform data cleansing, transformation, aggregations, etc.

Enables you to build resilient data flows in a code free environment

Enable you to focus on building business logic and data transformation

Underlying infrastructure is provisioned automatically with cloud scale via Spark execution

Mapping Data Flow



174

Lesson 04: Integrate Azure Data Factory with Azure Databricks

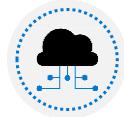


175

Lesson objectives



Use Azure Data Factory (ADF) to ingest data and create an ADF pipeline



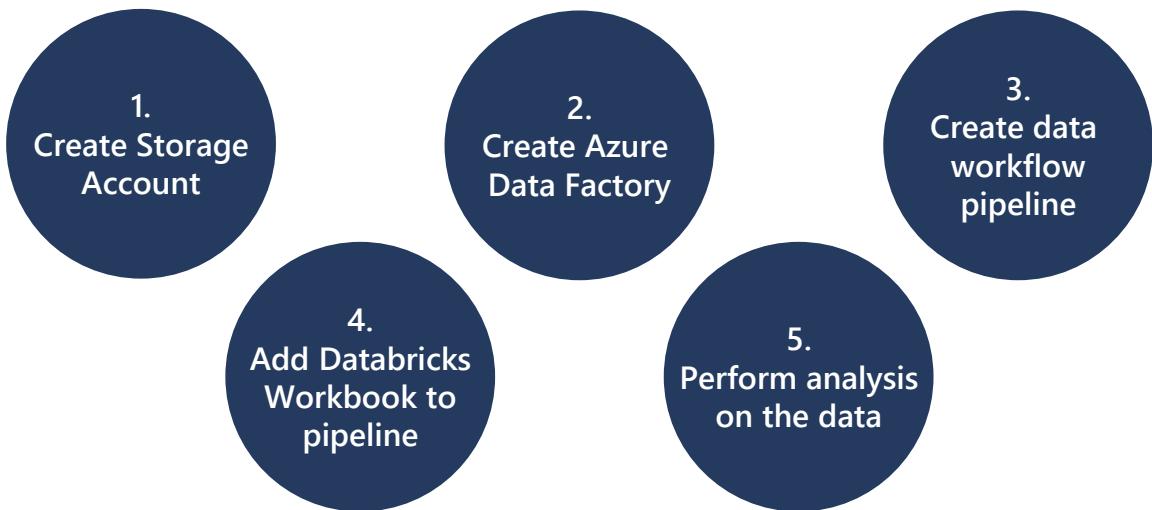
Create Azure Storage account and the Azure Data Factory instance



Use ADF to orchestrate data transformations using a Databricks Notebook activity

176

Working with documents programmatically



177

Create Azure Storage account and the Azure Data Factory instance

Create storage account

This screenshot shows the 'Create storage account' wizard. It has tabs for Basics, Advanced, Tags, and Review + create. The Basics tab displays information about Azure Storage and its features. The Project Details section allows selecting a subscription and resource group. The Instance Details section includes fields for Storage account name (e.g., 'contoso'), Location (West Europe), Performance (Standard), Account kind (StorageV2), Replication (Read-access geo-redundant storage (RA-GRS)), and Access tier (Hot). At the bottom, there are 'Review + create' and 'Next : Advanced >' buttons.

New data factory

This screenshot shows the 'New data factory' wizard. It requires entering a Name (e.g., 'v2'), Version (V2), Subscription (chtestao), and Resource Group (Select existing... or Create new). The Location is set to South Central US. Under the GIT tab, 'Enable GIT' is checked, and fields for GIT URL, Repo name, Branch Name, and Root folder are provided. A 'Create' button is at the bottom.

178

Use ADF to orchestrate data transformations using a Databricks Notebook activity

Data Transformation via Azure Data Factory

As you saw at the end of the previous lesson, different cities use different field names and values to indicate crimes, dates, etc. within their crime data.

For example:

- Some cities use the value "HOMICIDE", "CRIMINAL HOMICIDE" or "MURDER".
- In the New York data, the column is named `offenseDescription` while in the Boston data, the column is named `OFFENSE_CODE_GROUP`.
- In the New York data, the date of the event is in the `reportDate`, while in the Boston data, there is a single column named `MONTH`.

In the case of New York and Boston, here are the unique characteristics of each data set:

	Offense-Column	Offense-Value	Reported-Column	Reported-Data Type
New York	<code>offenseDescription</code>	starts with "murder" or "homicide"	<code>reportDate</code>	<code>timestamp</code>
Boston	<code>OFFENSE_CODE_GROUP</code>	"Homicide"	<code>MONTH</code>	<code>integer</code>

In this notebook, we will use an ADF Databricks Notebooks activity to perform transformations on and extract homicide statistics from the crime data being processed.

In this lesson you:

- Create Databricks Access Token.
- Add Databricks Notebook activity to pipeline.
- Connect Copy Activities to Notebook Activity.
- Publish the updated pipeline.
- Trigger and Monitor the pipeline run.
- Verify transformations of data by looking at the generated table in Databricks.
- Perform a simple aggregation of the data.

179

Lab: Orchestrating data movement with Azure Data Factory



180

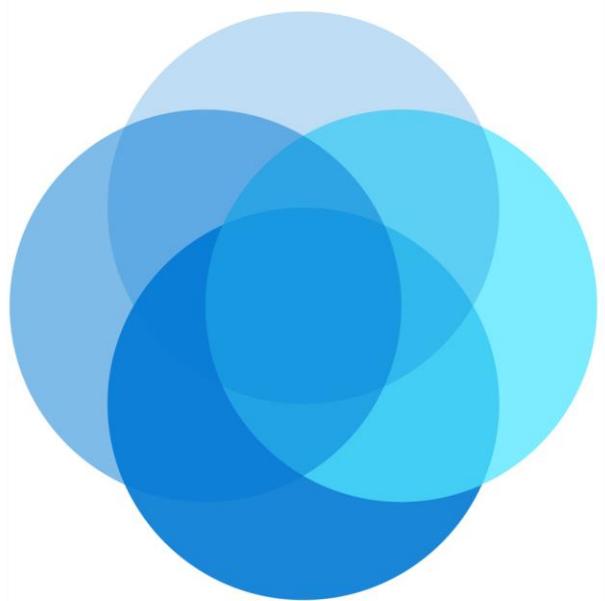


© Copyright Microsoft Corporation. All rights reserved.

181



Module 08: Securing Azure Data Platforms



182

Agenda



Lesson 01 – An introduction to security



Lesson 02 – Key security components



Lesson 03 – Securing storage accounts and Data Lake Storage



Lesson 04 – Securing data stores



Lesson 05 – Securing streaming data

183

Lesson 01: An Introduction to security

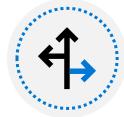


184

Lesson objectives



Shared security responsibility



A layered approach to security



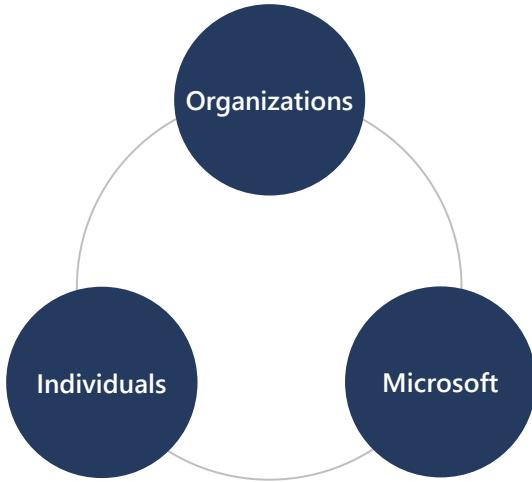
The Azure security center



Azure Government

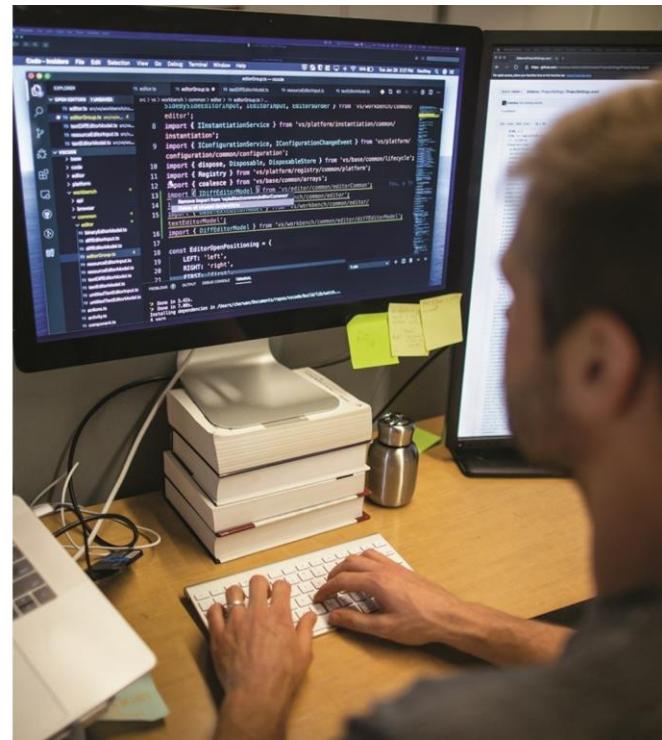
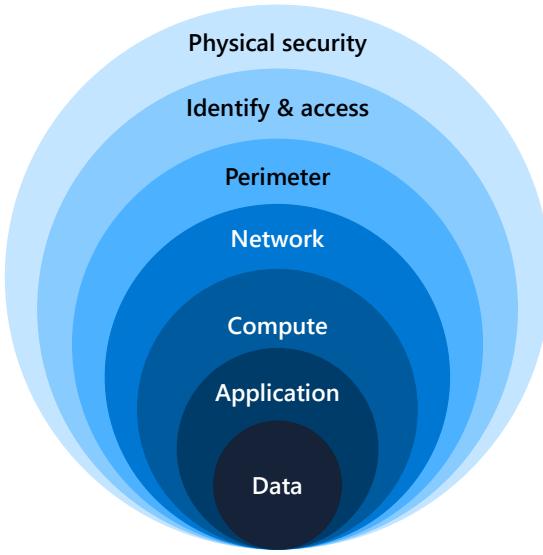
185

Shared security responsibility



186

A layered approach to security



187

Azure security center

Azure Security Center

Gain unmatched hybrid security management and threat protection

[Turn on Security Center >](#)

Not yet subscribed to Azure? [Start free >](#)

Pricing > Documentation > Updates > Training >

Turn on protection you need

Microsoft uses a wide variety of physical, infrastructure, and additional actions you need to take to help safeguard your security posture and protect against threats.

Use for incident response:

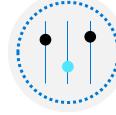
You can use Security Center during the detection, assessment, and diagnosis of security at various stages

Use to enhance security:

Reduce the chances of a significant security event by configuring a security policy, and then implementing the recommendations provided by Azure Security Center

188

Azure Government

-  Modernize Government services
-  Provide a platform of agility
-  Advanced Government mission
-  Physically separate from Azure

189

Lesson 02: Key security components



190

Lesson objectives



Network security



Identity and access management



Encryption capabilities built into Azure



Azure threat protection

191

Network security

Securing your network from attacks and unauthorized access is an important part of any architecture

Internet protection	Firewalls	DDoS protection	Network security groups
Assess the resources that are internet-facing, and to only allow inbound and outbound communication where necessary. Make sure you identify all resources that are allowing inbound network traffic of any type	To provide inbound protection at the perimeter, there are several choices: <ul style="list-style-type: none"> • Azure Firewall • Azure Application Gateway • Azure Storage Firewall 	The Azure DDoS Protection service protects your Azure applications by scrubbing traffic at the Azure network edge before it can impact your service's availability	Network Security Groups allow you to filter network traffic to and from Azure resources in an Azure virtual network. An NSG can contain multiple inbound and outbound security rules

192

Identity and access

Authentication

This is the process of establishing the identity of a person or service looking to access a resource. Azure Active Directory is a cloud-based identity service that provides this capability.

Authorization

This is the process of establishing what level of access an authenticated person or service has. It specifies what data they're allowed to access and what they can do with it. Azure Active Directory also provides this capability.

Azure Active Directory features

Single sign-on

Enables users to remember only one ID and one password to access multiple applications

Apps & device management

You can manage your cloud and on-premises apps and devices and the access to your organization's resources

Identity services

Manage Business-to-business (B2B) identity services and Business-to-Customer (B2C) identity services

193

Encryption

Encryption at rest

Data at rest is the data that has been stored on a physical medium. This could be data stored on the disk of a server, data stored in a database, or data stored in a storage account.

Encryption in transit

Data in transit is the data actively moving from one location to another, such as across the internet or through a private network. Secure transfer can be handled by several different layers.

Encryption on Azure

Raw encryption

Enables the encryption of:

- Azure Storage
- V.M. Disks
- Disk Encryption

Database encryption

Enables the encryption of databases using:

- Transparent Data Encryption

Encrypting secrets

Azure Key Vault is a centralized cloud service for storing your application secrets

194

Azure threat protection

The screenshot shows a timeline of threat activities for the domain contoso-corp. The events listed are:

- 4:04 PM Today**: Honeypot activity. Logged in to 2 computers via Contoso-DC.
- 3:21 PM Jan 22, 2018**: Remote execution attempt detected. Attempted remote execution of one or more WMI methods by AdminUser.
- 3:06 PM Jan 22, 2018**: Suspicious service creation. Administrator created 10 services in order to execute potentially malicious commands on Contoso-UC.
- 3:03 PM Jan 22, 2018**: Brute force attack using LDAP simple bind. 200 password guess attempts were made on 2 accounts from ALICE-DESKTOP; 2 account passwords were successfully guessed.
- 2:09 PM Jan 22, 2018**: Reconnaissance using account enumeration. Suspicious account enumeration activity using Kerberos protocol, originating from ALICE-DESKTOP, was detected. The attacker performed a total of 101 guess attempts for account names. 2 guess attempts matched existing account names in Active Directory.
- 12:38 PM Jan 21, 2018**: Malicious replication of directory services. Malicious replication requests were attempted by Alice Lübeli, from ALICE-DESKTOP against Contoso-DC.
- 11:59 AM Jan 21, 2018**: Reconnaissance using DNS. Suspicious DNS activity was observed, originating from ALICE-DESKTOP (which is not a DNS server) against Contoso-DC.

195

Lesson 03: Securing storage accounts and Data Lake Storage



196

Lesson objectives



Storage account security features



Explore the authentication options available to access data: Storage account key | Shared access signature



Control network access to the data



Managing encryption



Azure Data Lake Storage Gen II security features

197

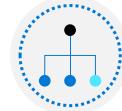
Storage account security features



Encryption
at rest



Encryption
in transit



Role based
access control



Auditing
access

198

Storage account keys

Home > Resource groups > cto_rg > ctoazureblob - Access keys

ctoazureblob - Access keys

Storage account

Search (Ctrl+)

Overview

Activity log

Access control (IAM)

Tags

Diagnose and solve problems

Events

Storage Explorer (preview)

Settings

Access keys (selected)

Geo-replication

CORS

Configuration

Encryption

Shared access signature

Use access keys to authenticate your applications when making requests to this Azure storage account. Store your access keys securely - for example, using Azure K Vault - and don't share them. We recommend regenerating your access keys regularly. You are provided two access keys so that you can maintain connections using one key while regenerating the other.

When you regenerate your access keys, you must update any Azure resources and applications that access this storage account to use the new keys. This action will interrupt access to disks from your virtual machines. [Learn more](#)

Storage account name: ctoazureblob

key1

Key: eU... Cg==

Connection string: Def... 3YQ...

key2

Key: NW... VpUgB5w==

Connection string: Def... ls6...

199

Shared access signatures

Home > Resource groups > cto_rg > ctoazureblob - Shared access signature

ctoazureblob - Shared access signature

Storage account

Search (Ctrl+)

Overview

Activity log

Access control (IAM)

Tags

Diagnose and solve problems

Events

Storage Explorer (preview)

Settings

Access keys

Geo-replication

CORS

Configuration

Encryption

Shared access signature (selected)

Firewalls and virtual networks

Advanced Threat Protection

Static website

Properties

Locks

Export template

Blob service

A shared access signature (SAS) is a URI that grants restricted access rights to Azure Storage resources. You can provide a shared access signature to clients who should not be trusted with your storage account key but whom you wish to delegate access to certain storage account resources. By distributing a shared access signature URI to these clients, you grant them access to a resource for a specified period of time.

An account-level SAS can delegate access to multiple storage services (i.e. blob, file, queue, table). Note that stored access policies are currently not supported for an account-level SAS.

[Learn more](#)

Allowed services: Block File Queue Table

Allowed resource types: Service Container Object

Allowed permissions: Read Write Delete List Add Create Update Process

Start and expiry date/time:

Start: 2019-03-29 11:59:33

End: 2019-03-29 19:59:33

(UTC+00:00) -- Current Time Zone

Allowed IP addresses: for example, 168.1.5.65 or 168.1.5.65-168.1.5.70

Allowed protocols: HTTPS only HTTPS and HTTP

Signing key: key1

Generate SAS and connection string

200

Control network access to data

Firewalls and virtual networks

Save Discard Refresh

Firewall settings allowing access to storage services will remain in effect for up to a minute after saving updated settings restricting access.

Allow access from All networks Selected networks

Configure network security for your storage accounts. [Learn more](#).

Virtual networks

Secure your storage account with virtual networks. [+ Add existing virtual network](#) [+ Add new virtual network](#)

VIRTUAL NETWORK	SUBNET	ADDRESS RANGE	ENDPOINT STATUS	RESOURCE GROUP	SUBSCRIPTION
No network selected.					

Firewall

Add IP ranges to allow access from the internet or your on-premises networks. [Learn more](#).

Add your client IP address (86.184.235.180) [?](#)

ADDRESS RANGE

IP address or CIDR

Exceptions

Allow trusted Microsoft services to access this storage account. [?](#)

Allow read access to storage logging from any network

Allow read access to storage metrics from any network

201

Managing encryption

Databases stores information that is sensitive, such as physical addresses, email addresses, and phone numbers. The following can be used to protect this data

Transport Layer Security (TLS)	Azure SQL Database and Data Warehouse enforces Transport Layer Security (TLS) encryption at all times for all connections, which ensures all data is encrypted "in transit" between the database and the client
Transparent data encryption	Both Azure Data Warehouse and SQL Database protects your data at rest using transparent data encryption (TDE). TDE performs real-time encryption and decryption of the database, associated backups, and transaction log files at rest without requiring changes to the application
Application encryption	Data in transit is a method to prevent man-in-the-middle attacks. To encrypt data in transit, specify Encrypt=true in the connection string in your client applications

202

Azure Data Lake Storage Gen2 security features

Role based access control

POSIX compliant ACL

AAD OAuth 2.0 token

Azure services integration

Encryption

203

Lesson 04: Securing data stores



204

Lesson objectives



Control network access to your data stores using firewall rules



Control user access to your data stores using authentication and authorization



Dynamic data masking



Audit and monitor your Azure SQL Database for access violations

205

Control network access to your data stores using firewall rules

There are a number of ways you can control access to your Azure SQL Database or Data Warehouse over the network

Server-level firewall rules

These rules enable clients to access your **entire Azure SQL server**, that is, all the databases within the same logical server

Database level firewall rules

These rules allow access to an individual database on a logical server and are stored in the database itself. For database-level rules, only **IP address rules** can be configured

206

Control user access to your data stores using authentication and authorization

Authentication

SQL Database and Azure Synapse Analytics supports two types of authentication: SQL authentication and Azure Active Directory authentication

Authorization

Authorization is controlled by permissions granted directly to the user account and/or database role memberships. A database role is used to group permissions together to ease administration

207

Dynamic data masking

Masking rules

MASK NAME	MASK FUNCTION
You haven't created any masking rules.	
SQL users excluded from masking (administrators are always excluded) ⓘ	
<i>SQL users excluded from masking (administrators are always excluded)</i> ✓	

Recommended fields to mask

SCHEMA	TABLE	COLUMN	
SalesLT	Address	AddressID	Add mask
SalesLT	Address	AddressLine1	Add mask
SalesLT	Address	AddressLine2	Add mask
SalesLT	Customer	FirstName	Add mask
SalesLT	Customer	LastName	Add mask

[Load more](#)

208

Auditing and monitoring



209

Lesson 05: Securing streaming data



210

Lesson objectives



Understand Stream Analytics security



Understand Event Hub security

211

Stream Analytics security

Data in transit

Azure Stream Analytics encrypts all incoming and outgoing communications and supports Transport Layer Security v 1.2

Data at rest

Stream Analytics doesn't store the incoming data since all processing is done in-memory. Therefore, consider setting security for services such as Event Hubs or Internet of Things Hubs, or for data stores such as Cosmos DB

212

Event Hub security

Authentication

Authentication makes use of Shared Access Signatures and Event Publishers to ensure that only applications or devices with valid credentials are only allowed to send data to an Event Hub. Each client is assigned a token

Token management

Once the tokens have been created, each client is provisioned with its own unique token. If a token is stolen by an attacker, the attacker can impersonate the client whose token has been stolen. Adding a client to a blocked recipients list renders that client unusable

213

Lab: Securing Azure Data Platforms



214

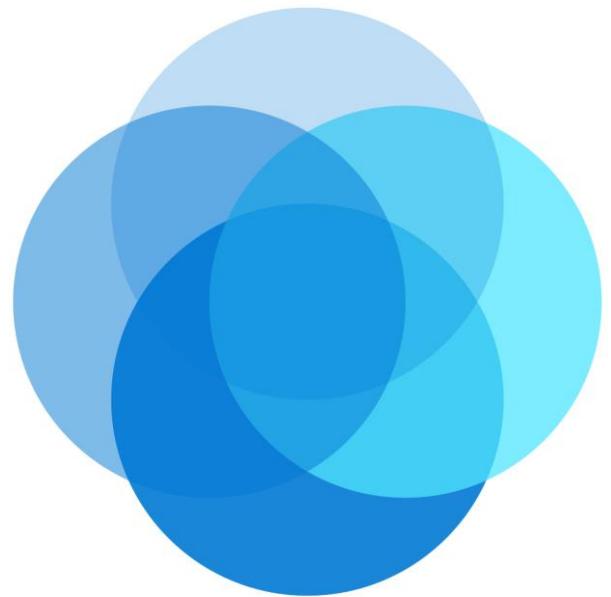


© Copyright Microsoft Corporation. All rights reserved.

215



Module 09: Monitoring and troubleshooting data storage and processing



216

Agenda



Lesson 01: General Azure monitoring capabilities



Lesson 02: Troubleshoot common data storage issues



Lesson 03: Troubleshoot common data processing issues



Lesson 04: Manage disaster recovery

217

Lesson 01: General Azure monitoring capabilities

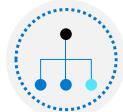


218

Lesson objectives



Azure Monitor



Monitoring the network



Diagnose and solve problems

219

Azure Monitor

Azure Monitor provides a holistic monitoring approach by collecting, analyzing, and acting on telemetry from both cloud and on-premises environments

Metric data

Provides quantifiable information about a system over time that enables you to observe the behavior of a system

Log data

Logs can be queried and even analyzed using Azure Monitor logs. In addition, this information is typically presented in the overview page of an Azure Resource in the Azure portal

Alerts

Alerts notify you of critical conditions and potentially take corrective automated actions based on triggers from metrics or logs

220

Monitoring the network

Azure Monitor logs within Azure monitor has the capability to monitor and measure network activity

Network performance monitor

Network Performance Monitor measures the performance and reachability of the networks that you have configured

Application gateway analytics

Application Gateway Analytics contains rich, out-of-the box views you can get insights into key scenarios, including:

- Monitor client and server errors
- Check requests per hour

221

Diagnose and solve issues

The screenshot shows the 'ctocdb - Diagnose and solve problems' blade in the Azure portal. The left sidebar lists various monitoring and diagnostic tools: Overview, Activity log, Access control (IAM), Tags, Diagnose and solve problems (selected), Quick start, Notifications, Data Explorer, Settings, Replicate data globally, Default consistency, Firewall and virtual networks, and CORS. The main content area is titled 'RESOURCE HEALTH' and shows a green 'Available' status with a note: 'There aren't any known problems affecting this Cosmos DB database account [More details](#)'. Below this is the 'RECENT ACTIVITY' section, which displays activity for the past 24 hours and links to 'Quick Insights' and 'See all activity'. The final section is 'SOLUTIONS TO COMMON PROBLEMS', which lists several troubleshooting items with dropdown arrows: 'My database is slow', 'My request unit (RU) charging is unclear', 'I need more storage/throughput', 'My queries are slow', 'MongoDB API Support', and 'Import MongoDB data into CosmosDB'.

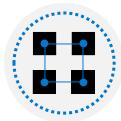
222

Lesson 02: Troubleshoot common data storage issues



223

Lesson objectives



Connectivity issues



Performance issues



Storage issues

224

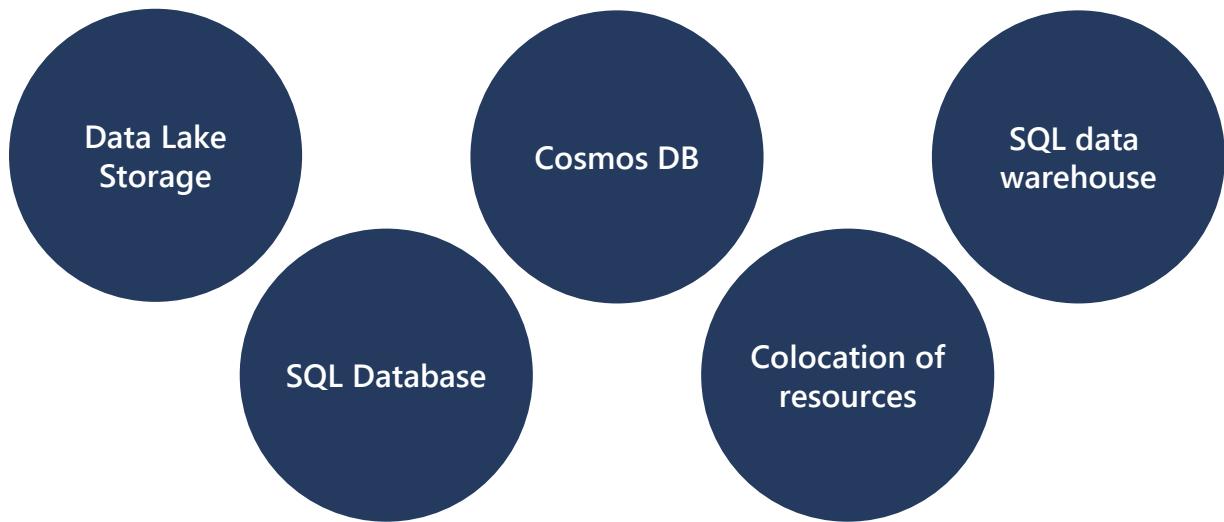
Connectivity issues

There are a range of issues that can impact connectivity issues, including:

Unable to connect to the data platform	Authentication failures	Cosmos DB Mongo DB API errors	SQL database failover
<p>The first area that you should check is the firewall configuration</p> <p>Test the connection by accessing it from a location external to your network</p> <p>Check maintenance schedules</p>	<p>The first check is to ensure that the username and password is correct</p> <p>Check the storage account keys and ensure that they match in the connection string</p>	<p>Mongo client drivers establishes more than one connection</p> <p>On the server side, connections which are idle for more than 30 minutes are automatically closed down</p> <p>Check for timeouts</p>	<p>Should you receive an "unable to connect" message (error code 40613) in the Azure SQL Database, this scenario commonly occurs when a database has been moved because of deployment, failover, or load balancing</p>

225

Performance issues



226

Storage issues

Consistency:

Consider the consistency levels of the following data stores that can impact data consistency:

- Cosmos DB
- SQL Data Warehouse
- SQL Database

Corruption:

Data corruption can occur on any of the data platforms for a variety of reasons. You should have an appropriate disaster recovery strategy

227

Lesson 03: Troubleshoot common data processing issues



228

Lesson objectives



Troubleshoot streaming data



Troubleshoot batch data loads



Troubleshoot Azure Data Factory

229

Troubleshoot streaming data

When using Stream Analytics, a Job encapsulates the Stream Analytic work and is made up of three components:

Job input

The job input contains a **Test Connection** button to validate that there is connectivity with the input. However, most errors associated with a job input is due to the malformed input data that is being ingested

Job query

A common issue associated with Stream Analytics query is the fact that the output produced is not expected. In this scenario it is best to check the query itself to ensure that there is no mistakes on the code there

Job output

As with the job input, there is a **Test Connection** button to validate that there is connectivity with the output, should there be no data appearing. You can also use the **Monitor** tab in Stream Analytics to troubleshoot issues

230

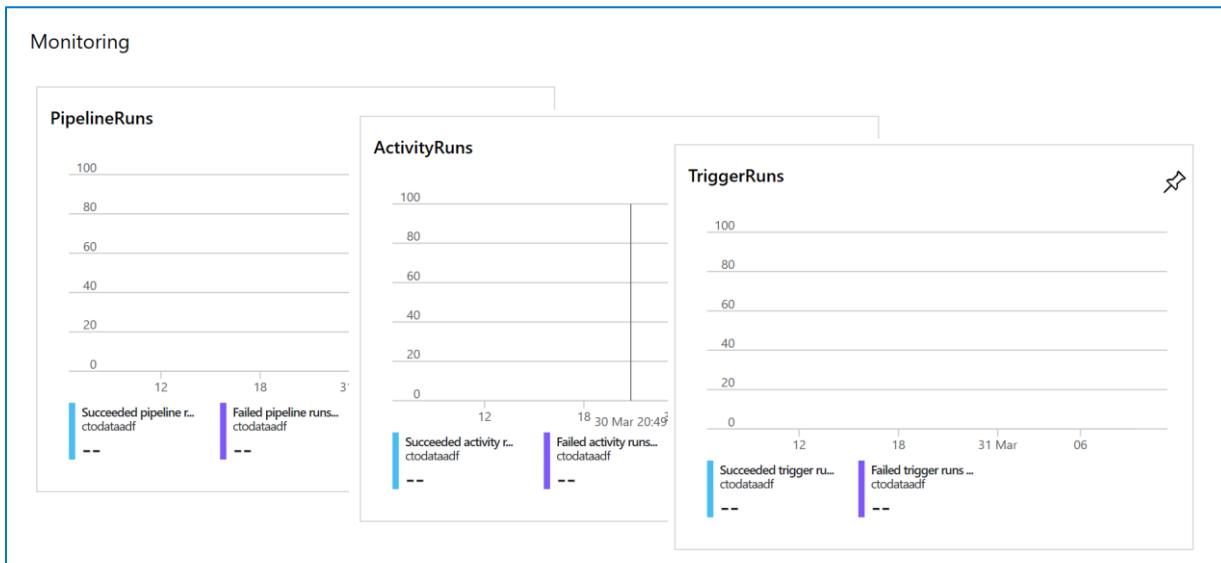
Troubleshoot batch data loads

When trying to resolve data load issues, it is first pragmatic to make the holistic checks on Azure, as well as the network checks and diagnose and solve issue check. After that, then check:

Azure Blob and Data Lake Store	SQL Data Warehouse	Cosmos DB	SQL Database
Notwithstanding network errors; occasionally, you can get timeout or throttling errors that can be a symptom of the availability of the storage accounts	<ul style="list-style-type: none"> Make sure you are always leveraging PolyBase Ensure CTAS statements are used to load data Break data down into multiple text files Consider DWU usage 	<ul style="list-style-type: none"> Check that you have provisioned enough RU's Review partitions and partitioning keys Check for client connection string settings 	<ul style="list-style-type: none"> Check that you have provisioned enough DTU's Review whether the database would benefit from elastic pools A wide range of tools can be used to troubleshoot SQL Database

231

Troubleshoot Azure Data Factory



232

Lesson 04: Managing disaster recovery



233

Lesson objectives



Data redundancy



Disaster recovery

234

Data redundancy

Data redundancy is the process of storing data in multiple locations to ensure that it is highly available

Azure Blob and Data Lake Store	SQL Data Warehouse	Cosmos DB	SQL Database
Locally redundant storage (LRS) Zone-redundant storage (ZRS) Geo-redundant storage (GRS) Read-access geo-redundant storage (RA-GRS)	SQL Data Warehouse performs a geo-backup once per day to a paired data center. The RPO for a geo-restore is 24 hours	Azure Cosmos DB is a globally distributed database service. You can configure your databases to be globally distributed and available in any of the Azure regions	Check that you have provisioned enough DTU's Review whether the database would benefit from elastic pools A wide range of tools can be used to troubleshoot SQL Database

235

Disaster Recovery

There should be processes that are involved in backing up or providing failover for databases in an Azure data platform technology. Depending on circumstances, there are numerous approaches that can be adopted

Azure Blob and Data Lake Store	SQL Data Warehouse	Cosmos DB	SQL Database
Supports account failover for geo-redundant storage accounts You can initiate the failover process for your storage account if the primary endpoint becomes unavailable	SQL Data Warehouse performs a geo-backup once per day to a paired data center Data warehouse snapshot feature that enables you to create a restore point to create a copy of the warehouse to a previous state	Takes a backup of your database every 4 hours and at any point of time Only the latest 2 backups are stored	Creates database backups that are kept between 7 and 35 days Uses Azure read-access geo-redundant storage (RA-GRS) to ensure that they preserved even if data center is unavailable

236

Lab: Monitoring and troubleshooting data storage and processing



237

 Microsoft Azure

© Copyright Microsoft Corporation. All rights reserved.

238